

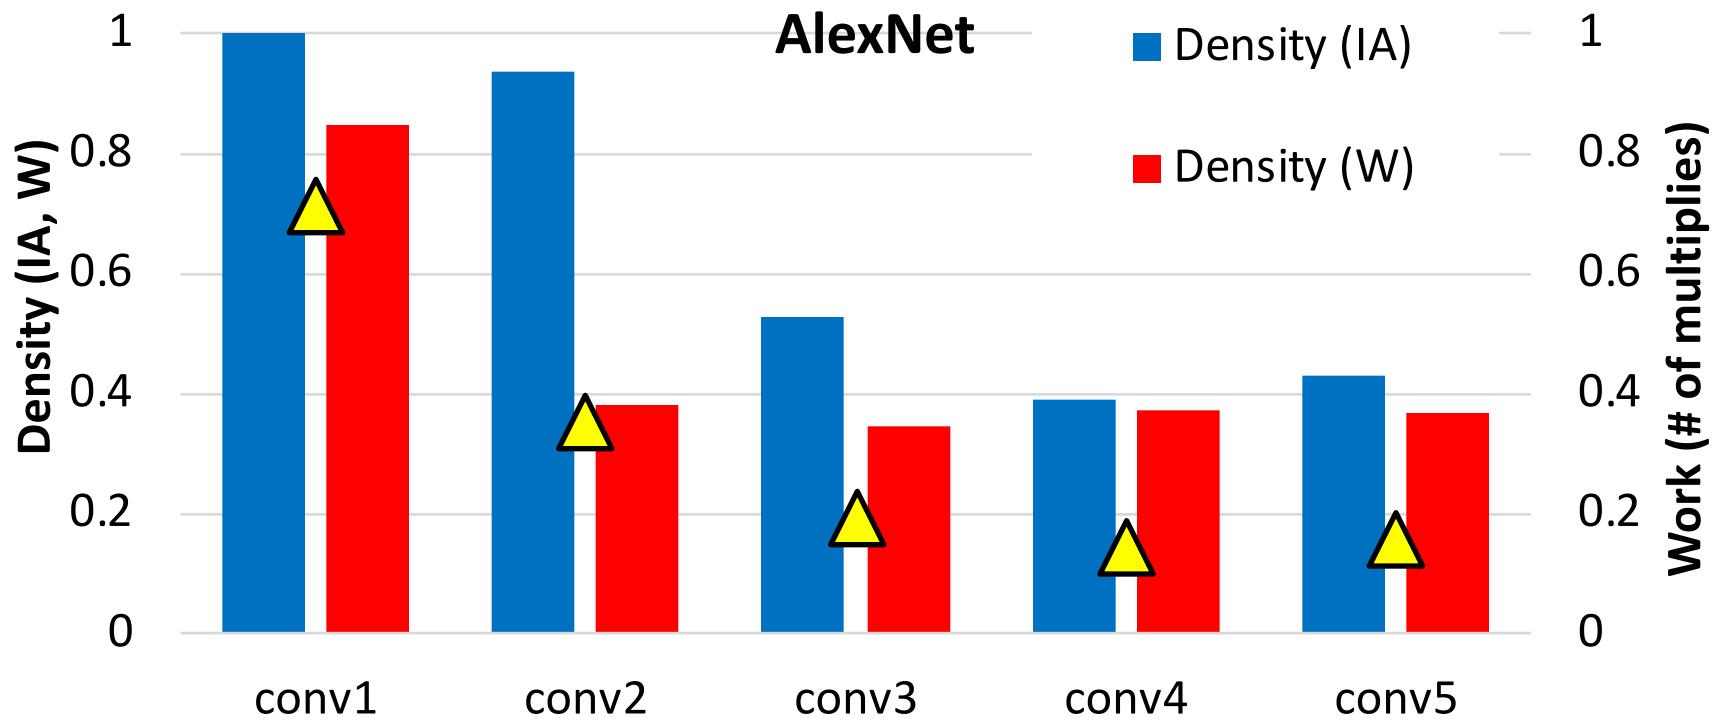
Sparse Architectures

ISCA Tutorial (2019)

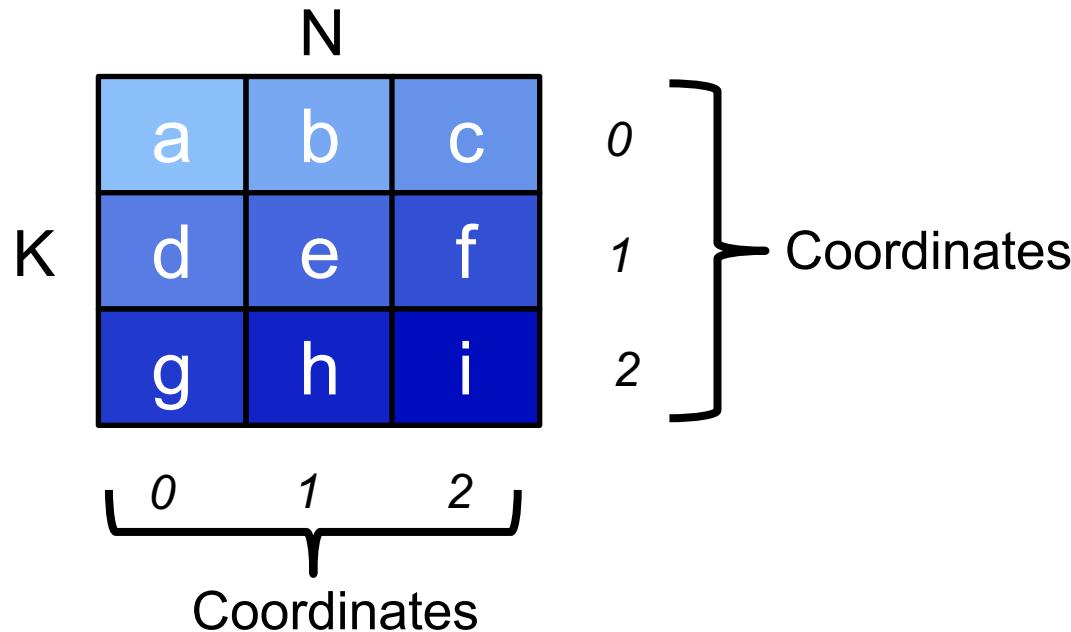
Website: <http://eyeriss.mit.edu/tutorial.html>

Motivation

- Leverage CNN sparsity to improve energy-efficiency

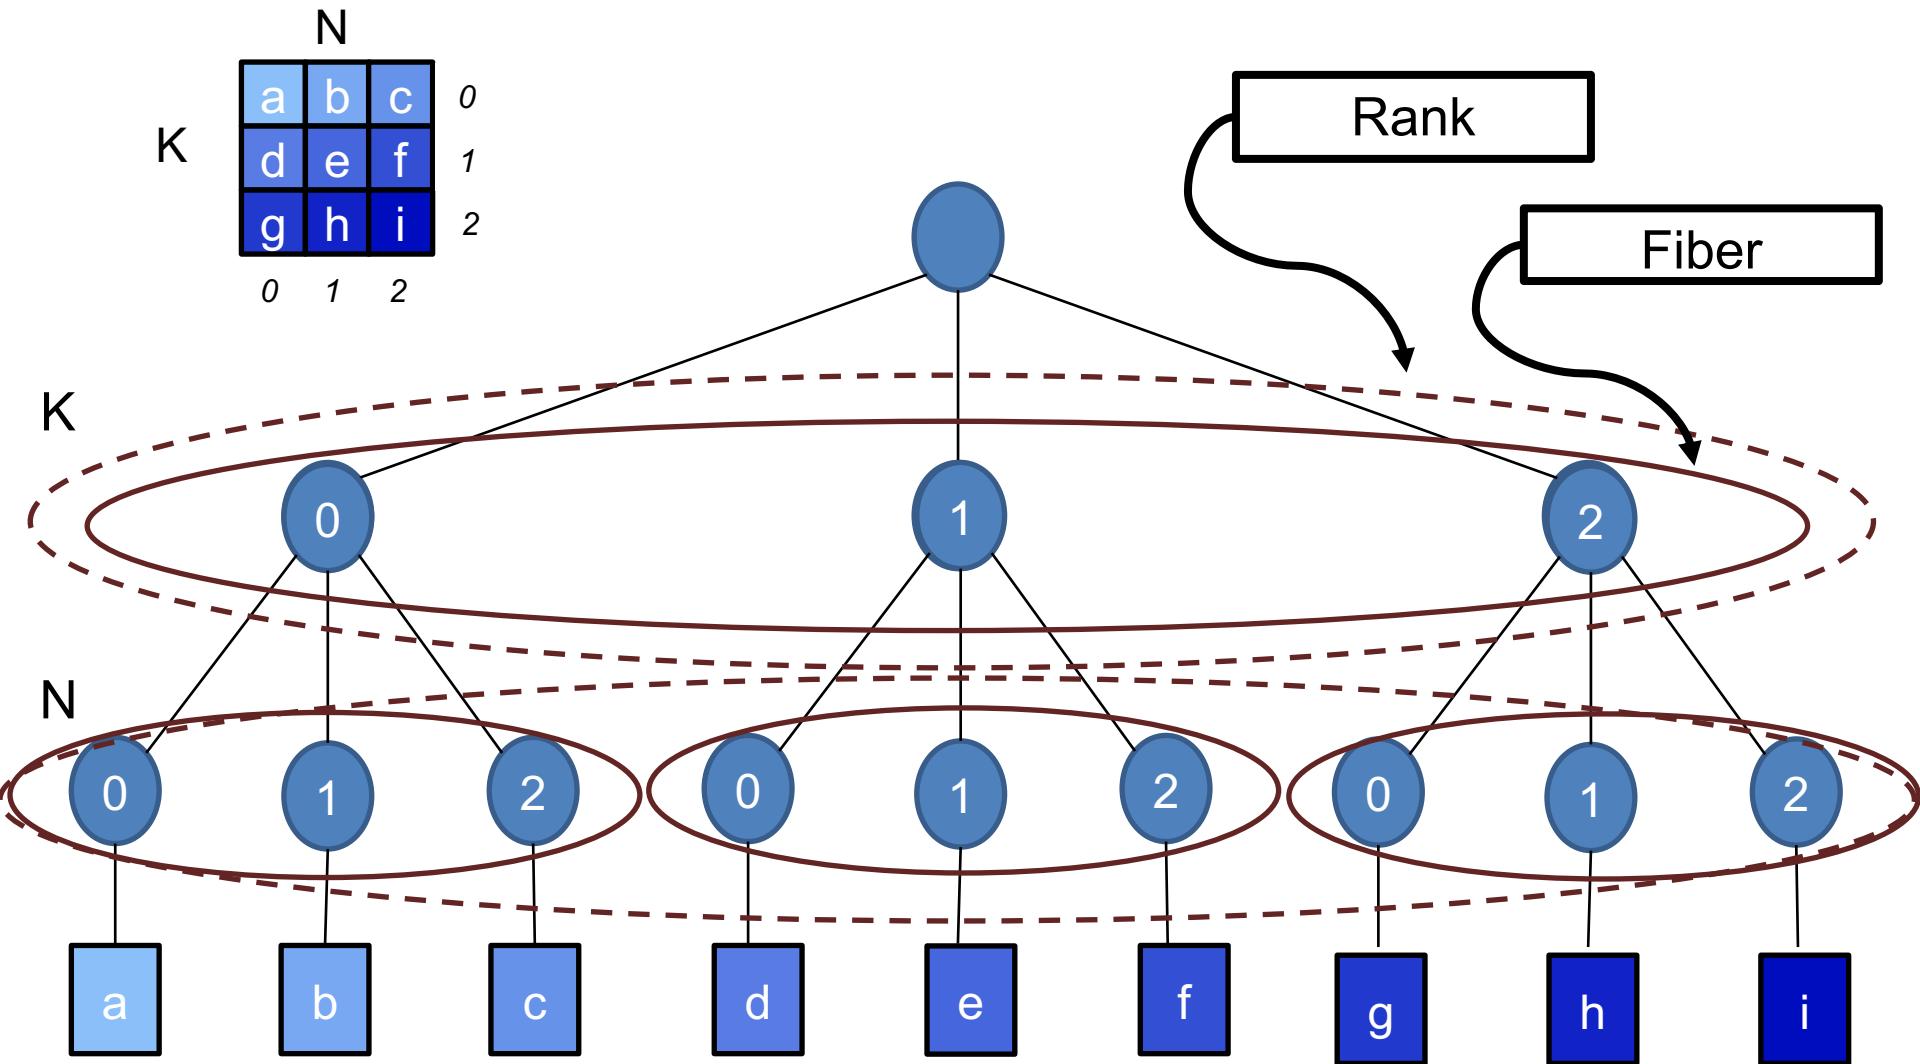


Tensor Data



- The elements of each “rank” (dimension) are identified by their “coordinates”, e.g., rank K has coordinates 0, 1, 2
- Each element of the tensor is identified by the tuple of coordinates from each of its ranks, i.e., a “point”.

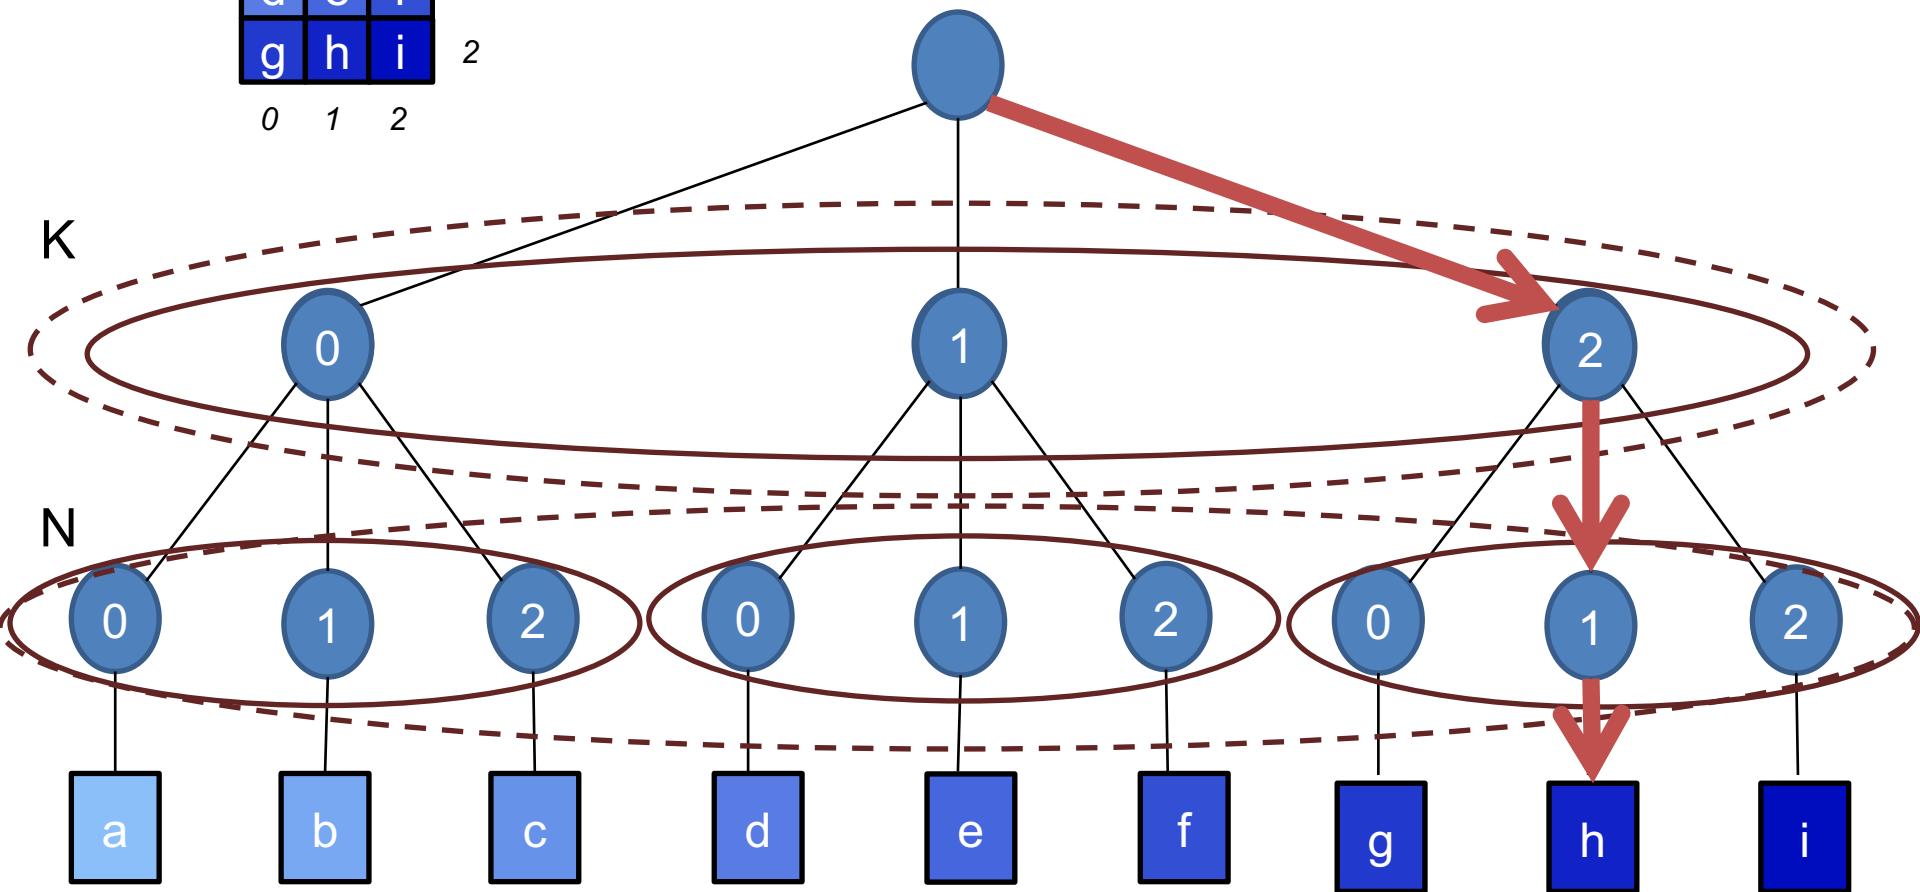
Tree-based Tensor Representation



Tree-based Tensor Representation

	N		
K	a	b	c
	d	e	f
	g	h	i
0	1	2	

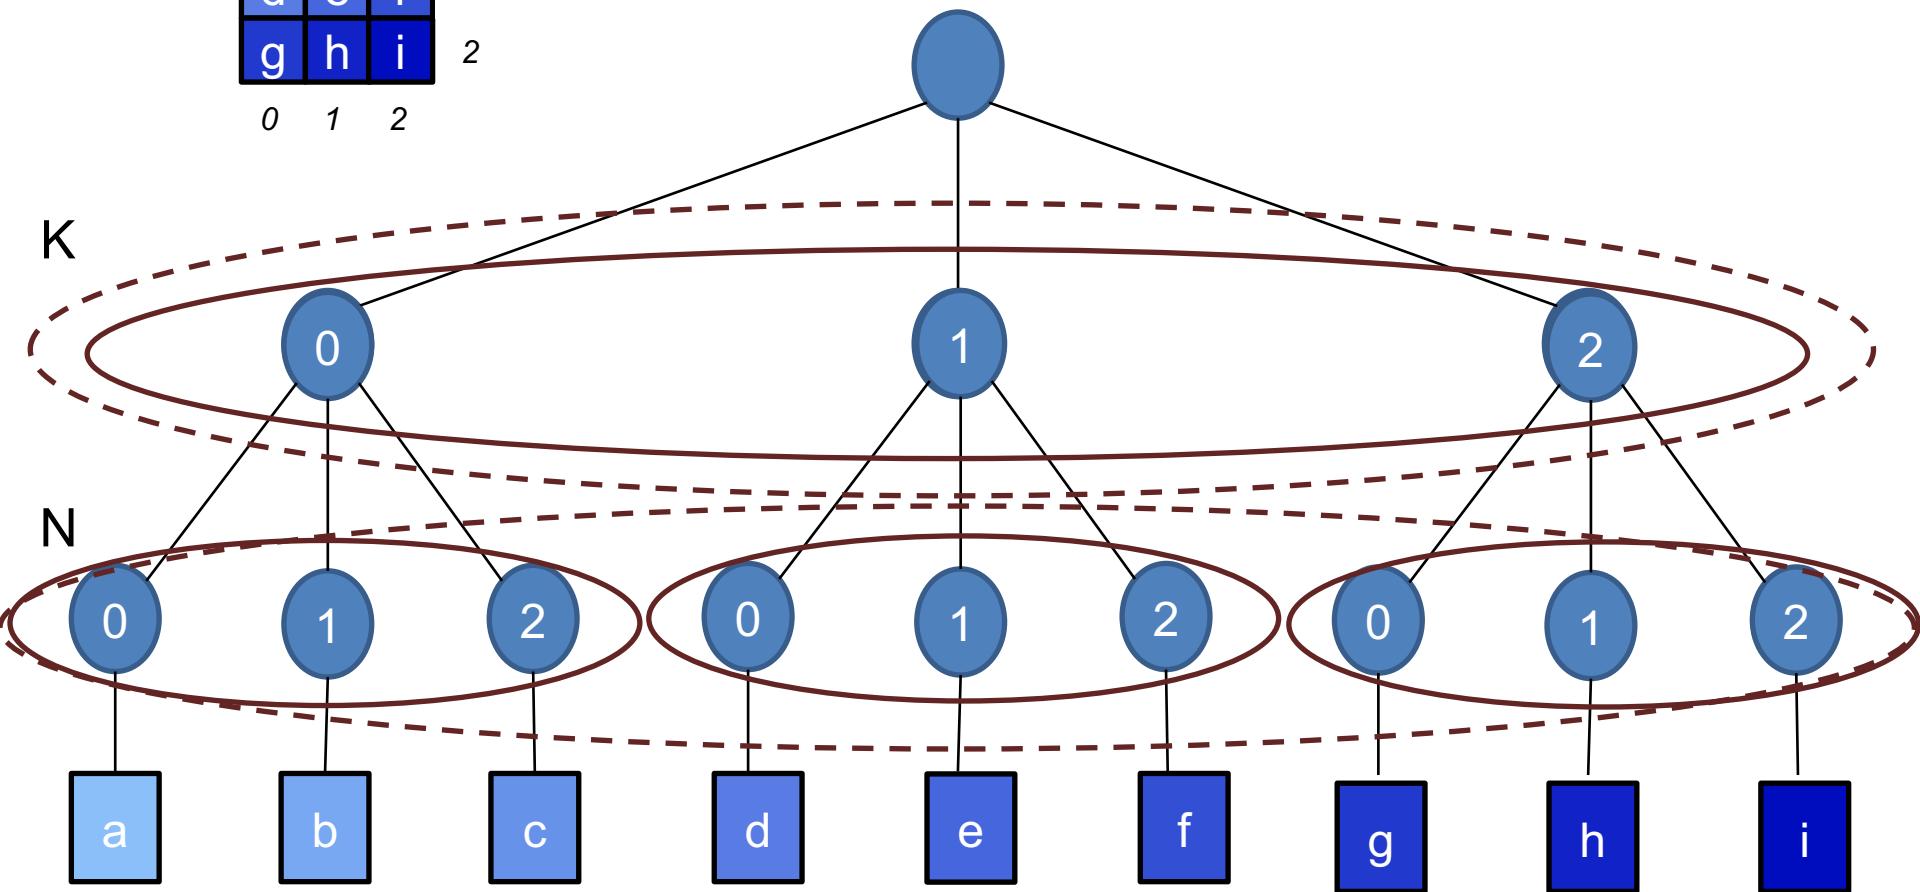
Finding point (2, 1)



Tree-based Tensor Representation

	N		
K	a	b	c
	d	e	f
	g	h	i
0	1	2	

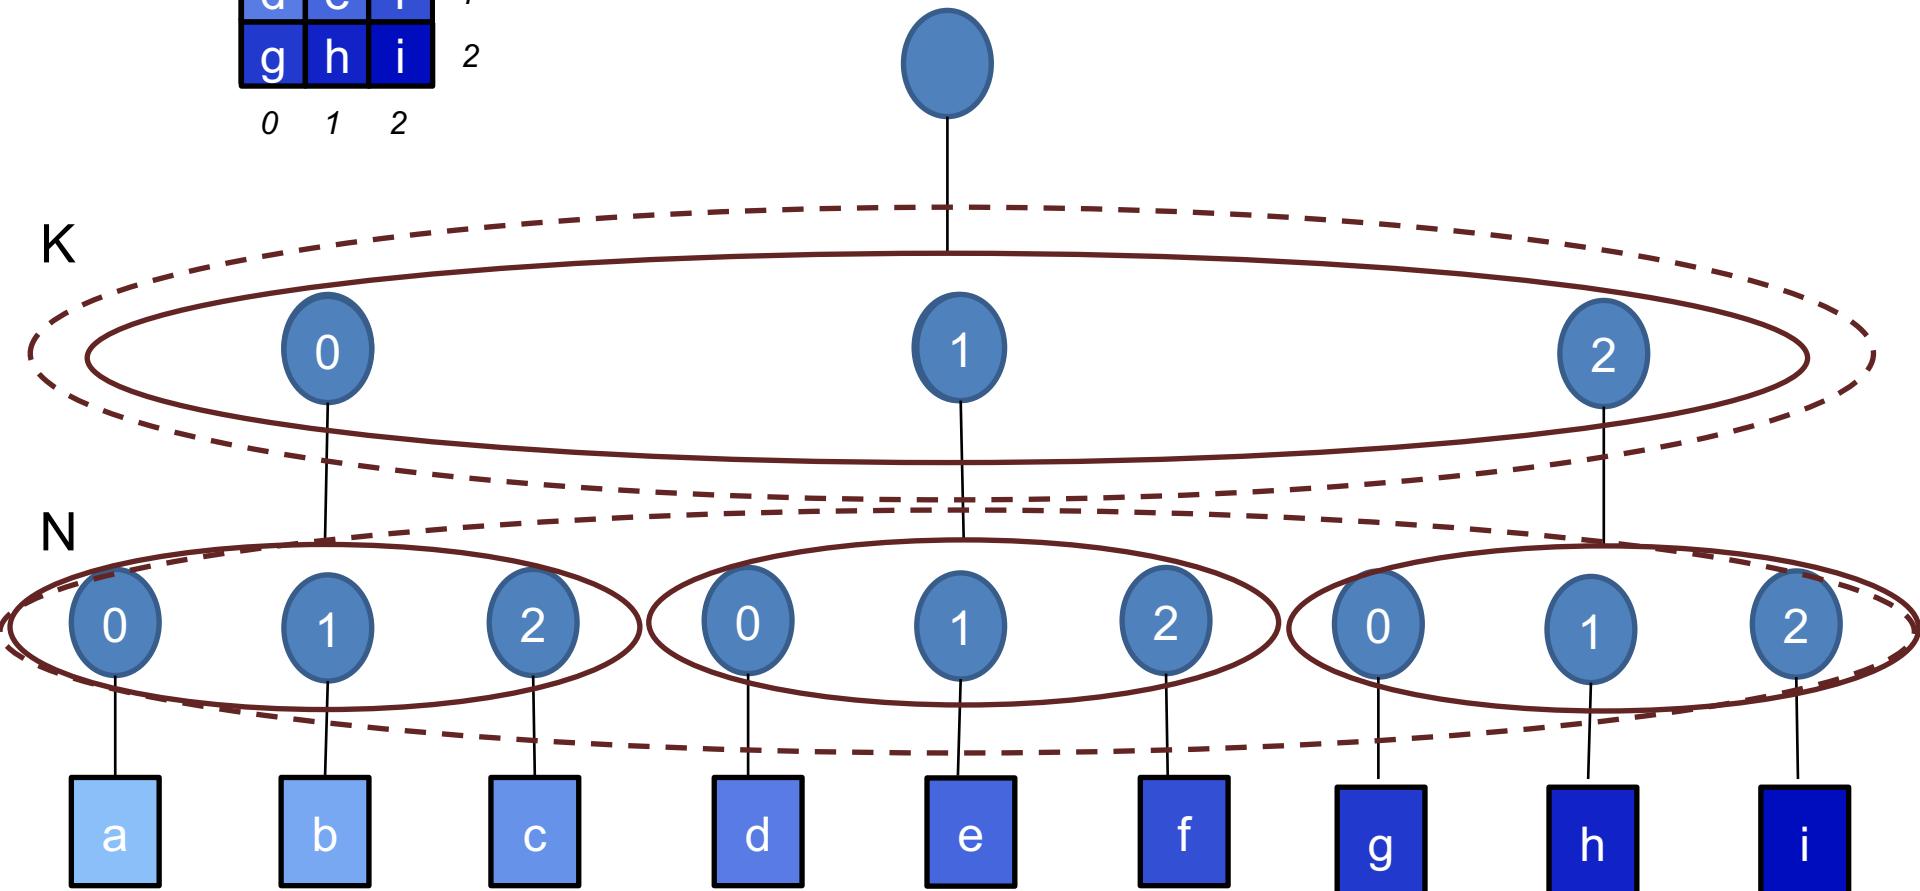
Each coordinate references a fiber



Rank-based Tensor Representation

	N		
K	a	b	c
	d	e	f
	g	h	i
0	1	2	

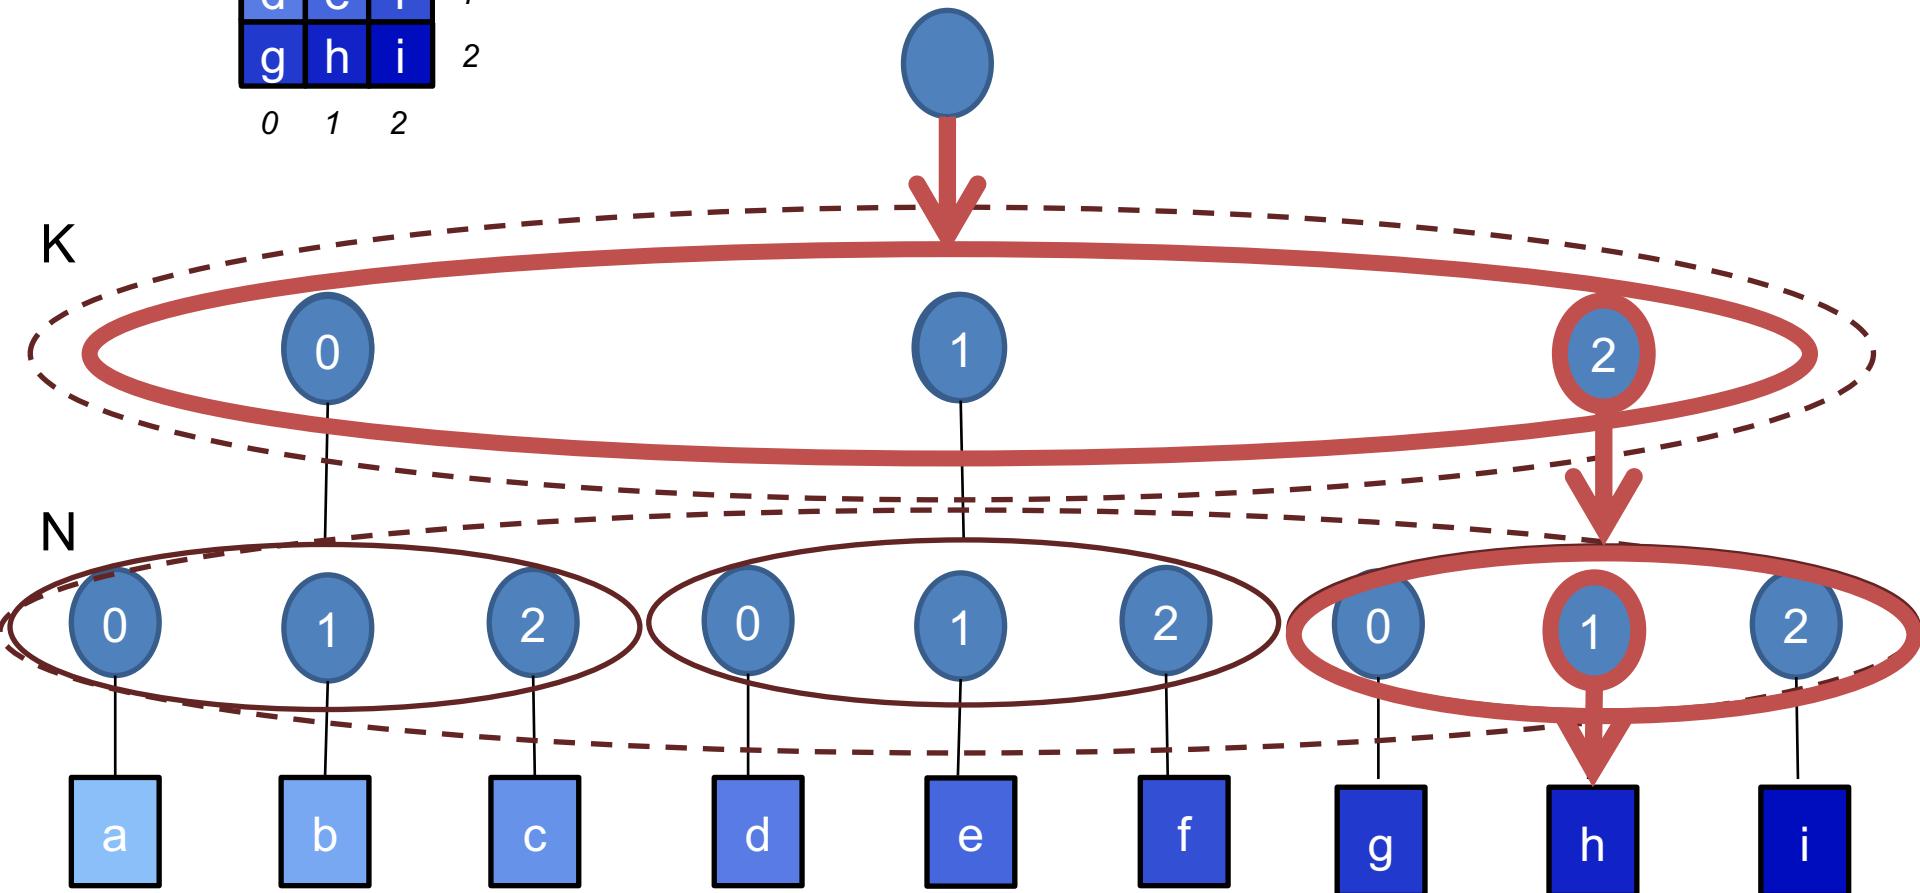
Each coordinate references a fiber



Rank-based Tensor Representation

	N		
K	a	b	c
	d	e	f
	g	h	i
0	1	2	

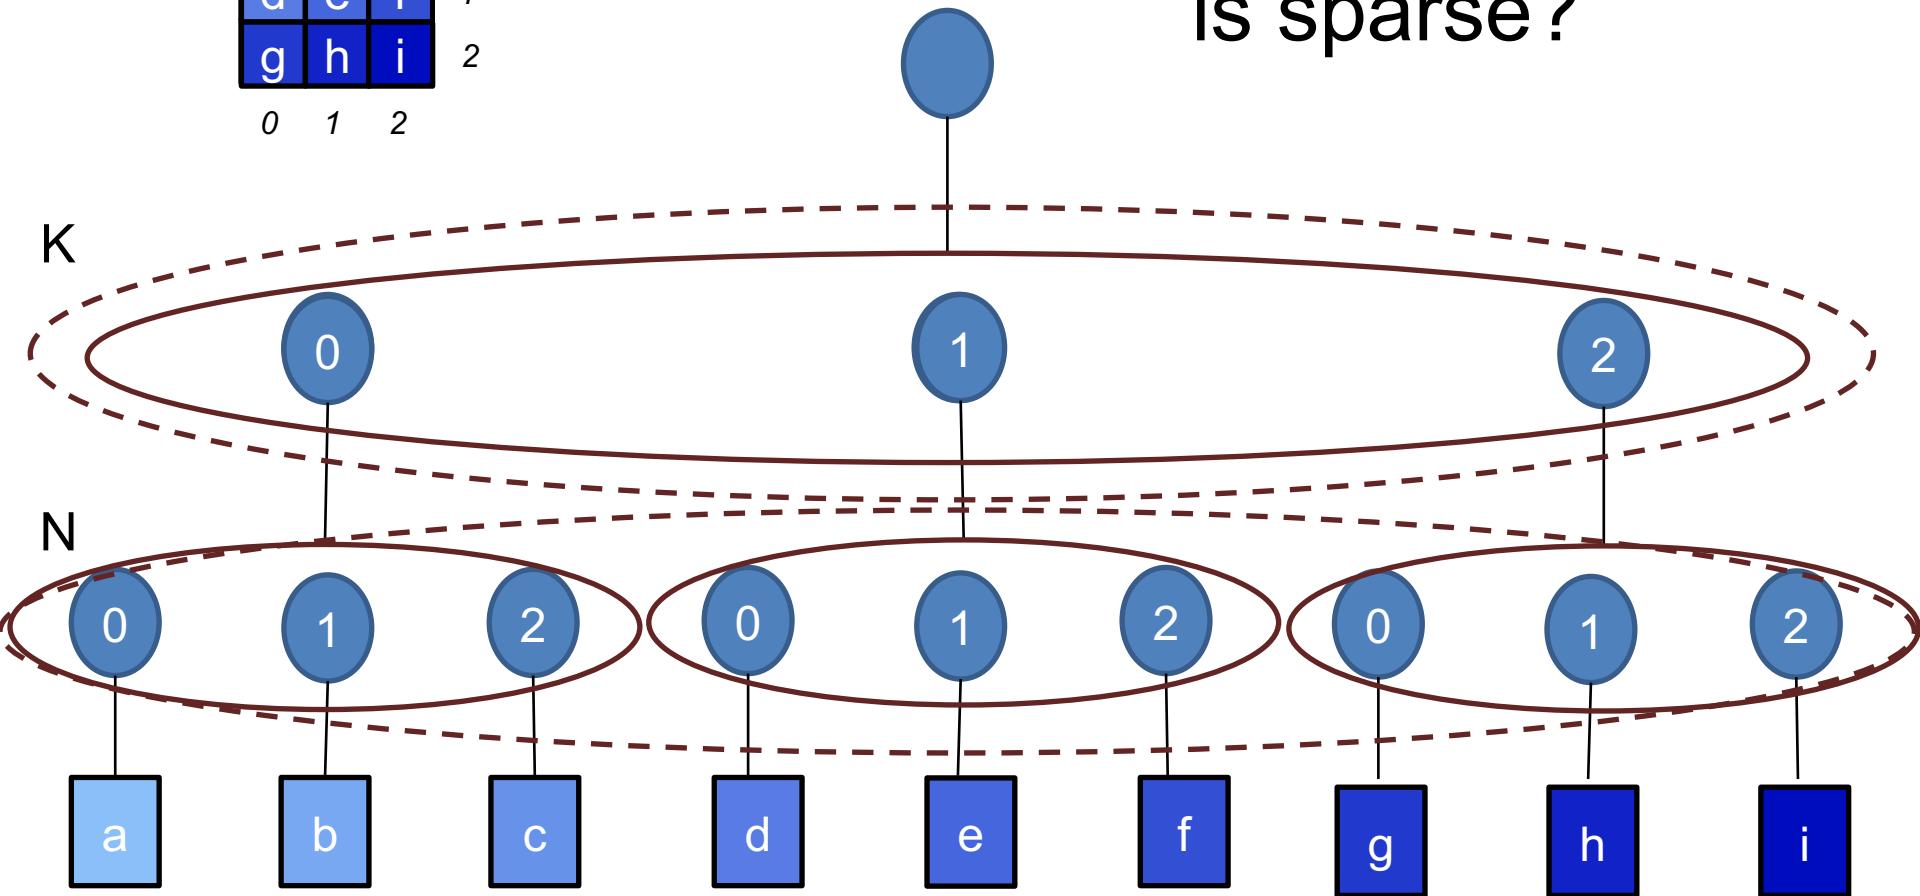
Finding point (2, 1)



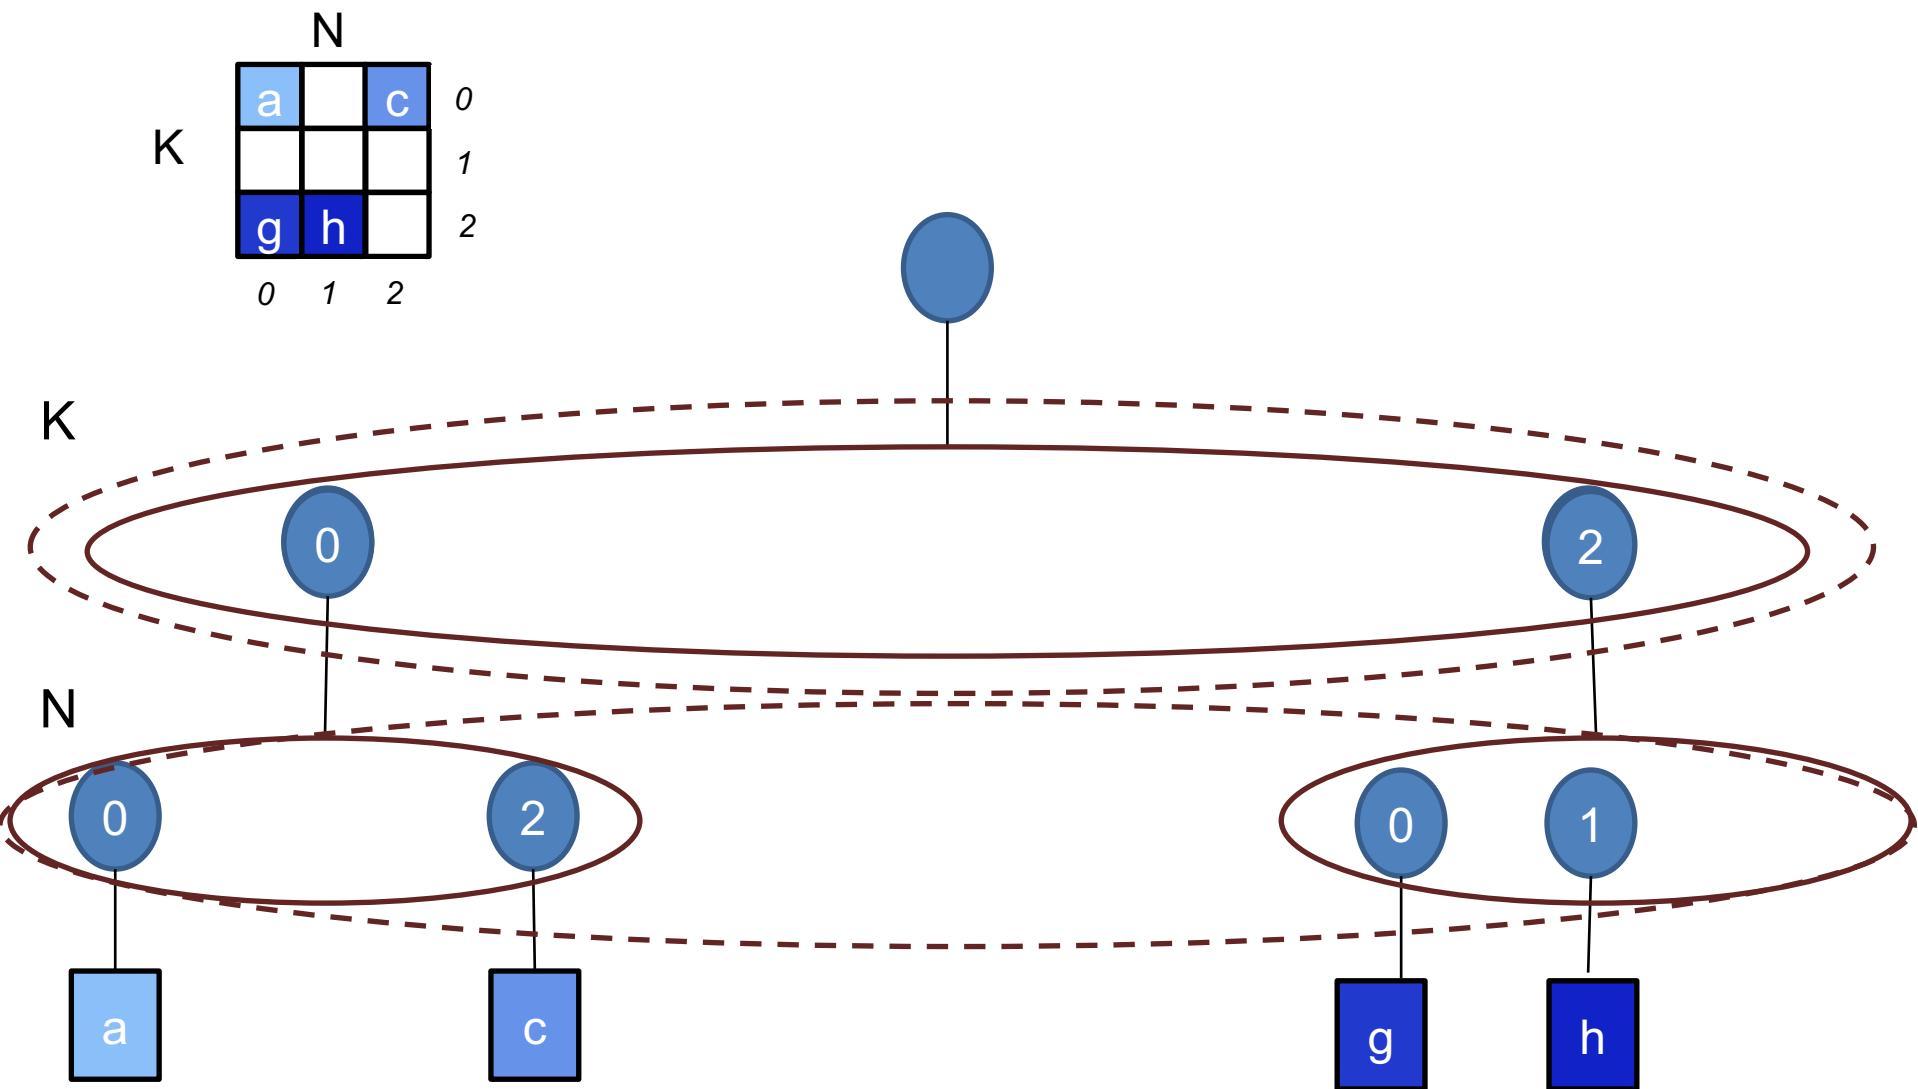
Rank-based Tensor Representation

	N
K	a b c
	d e f
	g h i
0	0 1 2

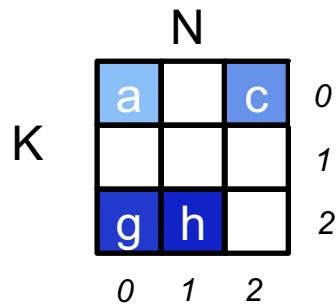
What if tensor
is sparse?



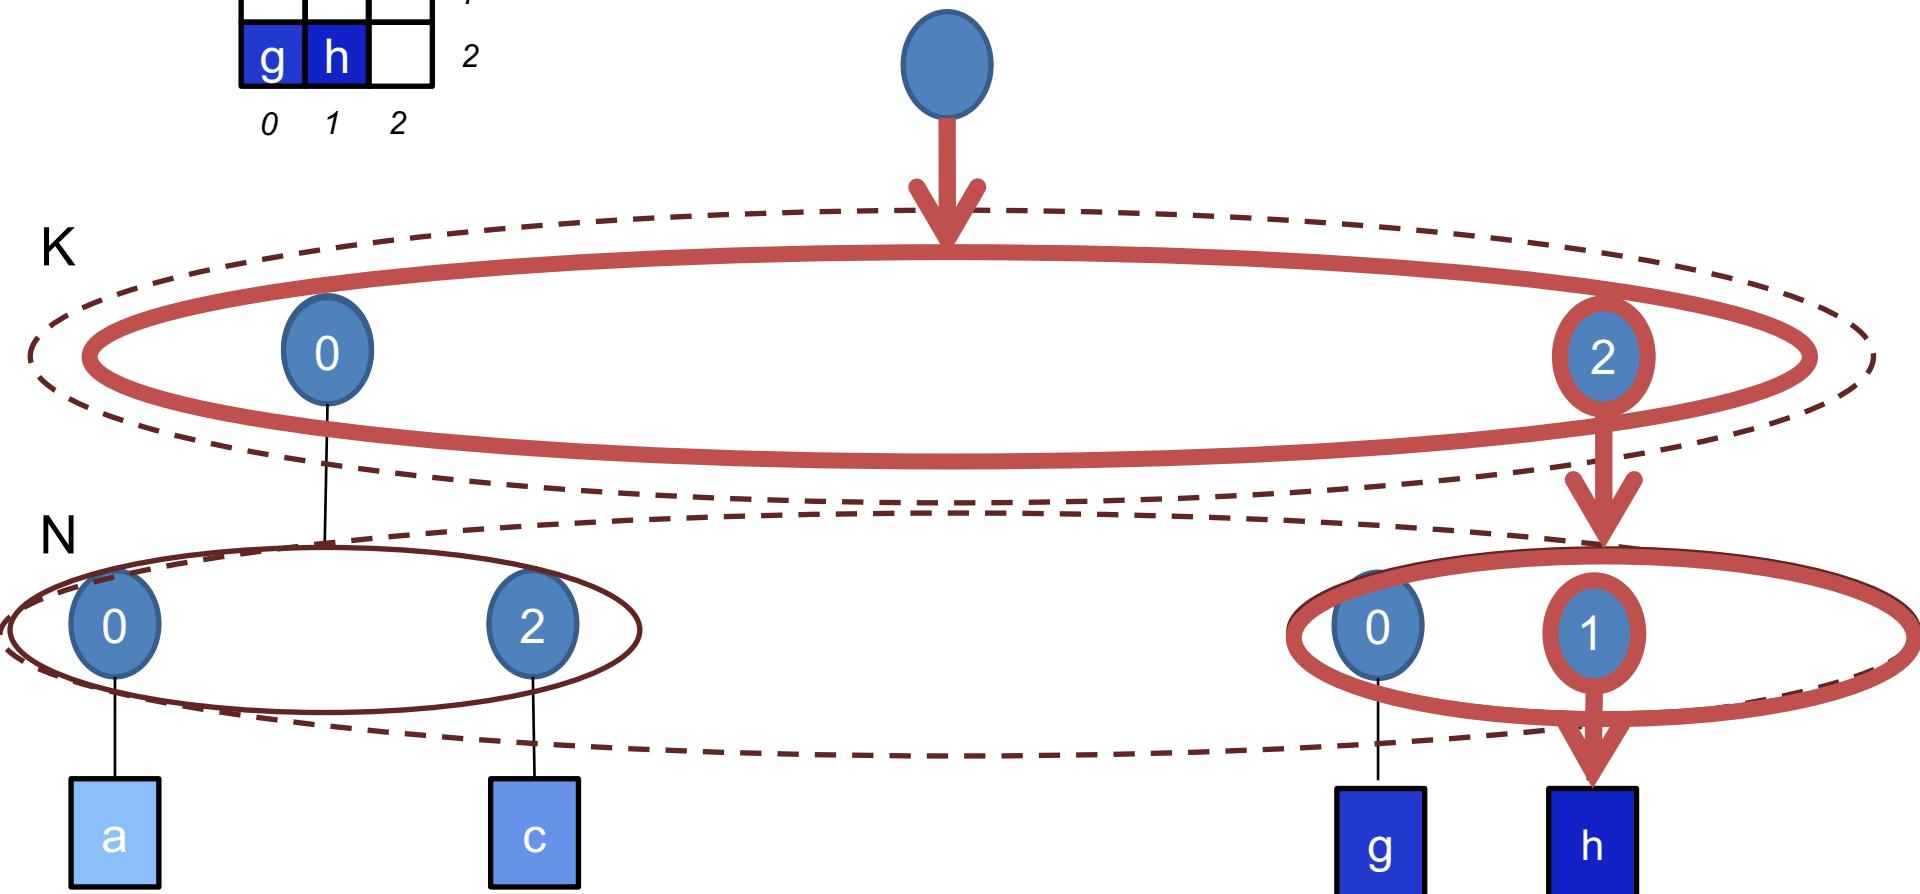
Sparse Tensor Representation



Sparse Tensor Representation

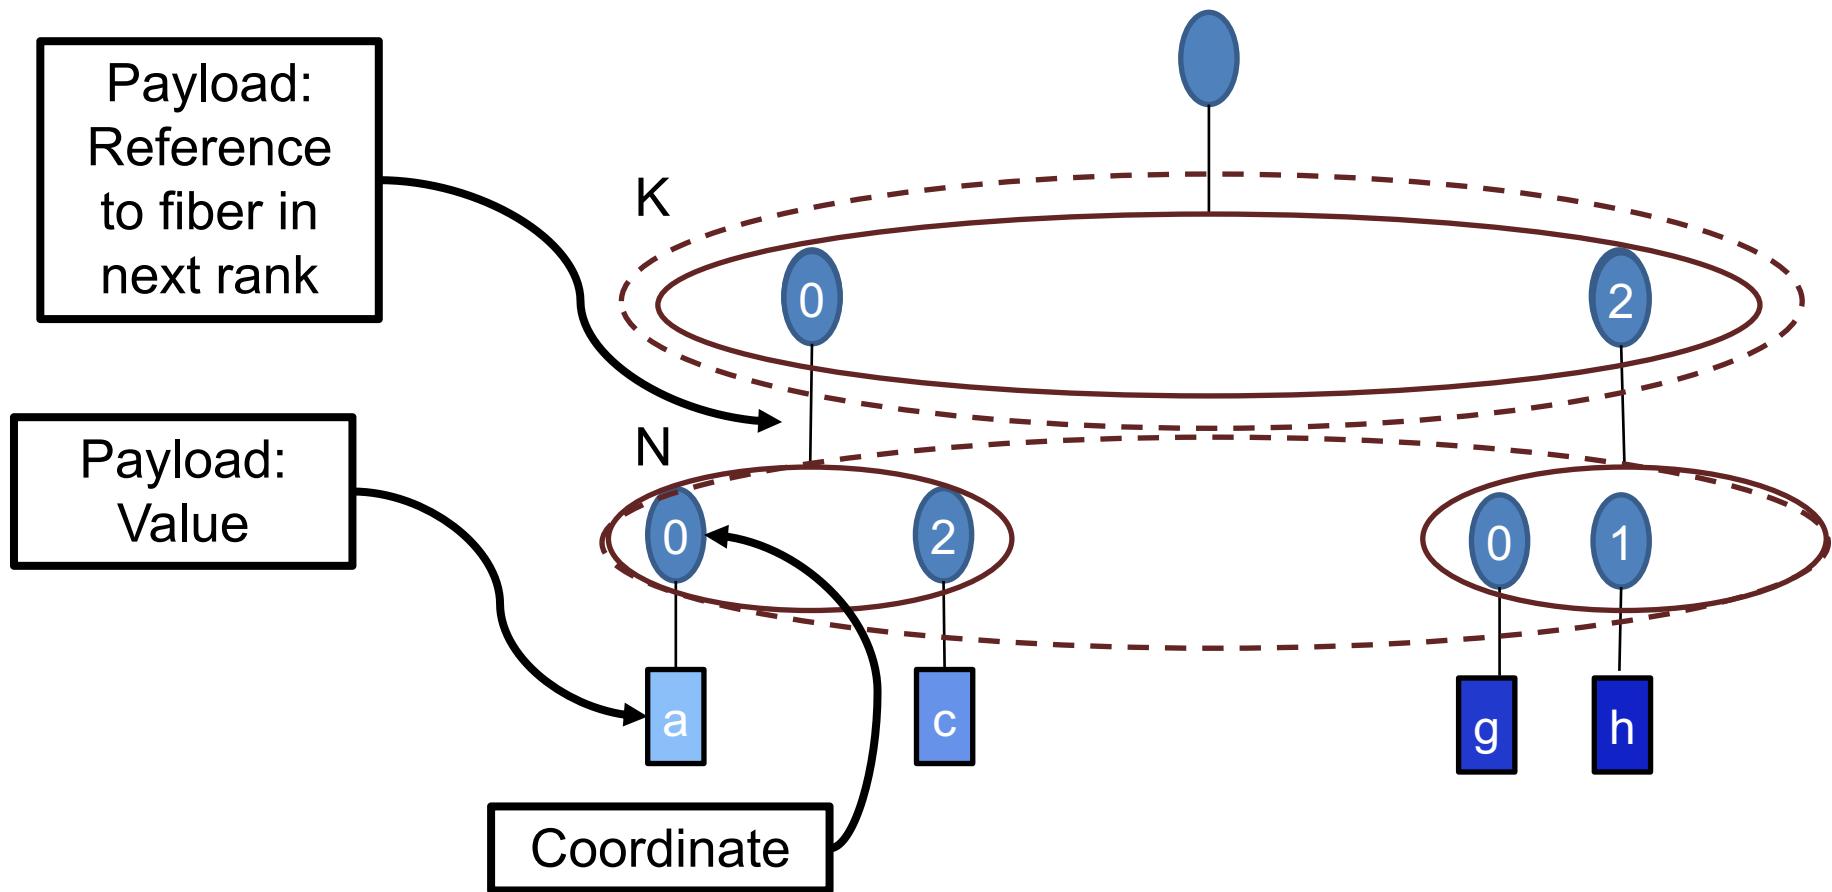


Finding point (2, 1)



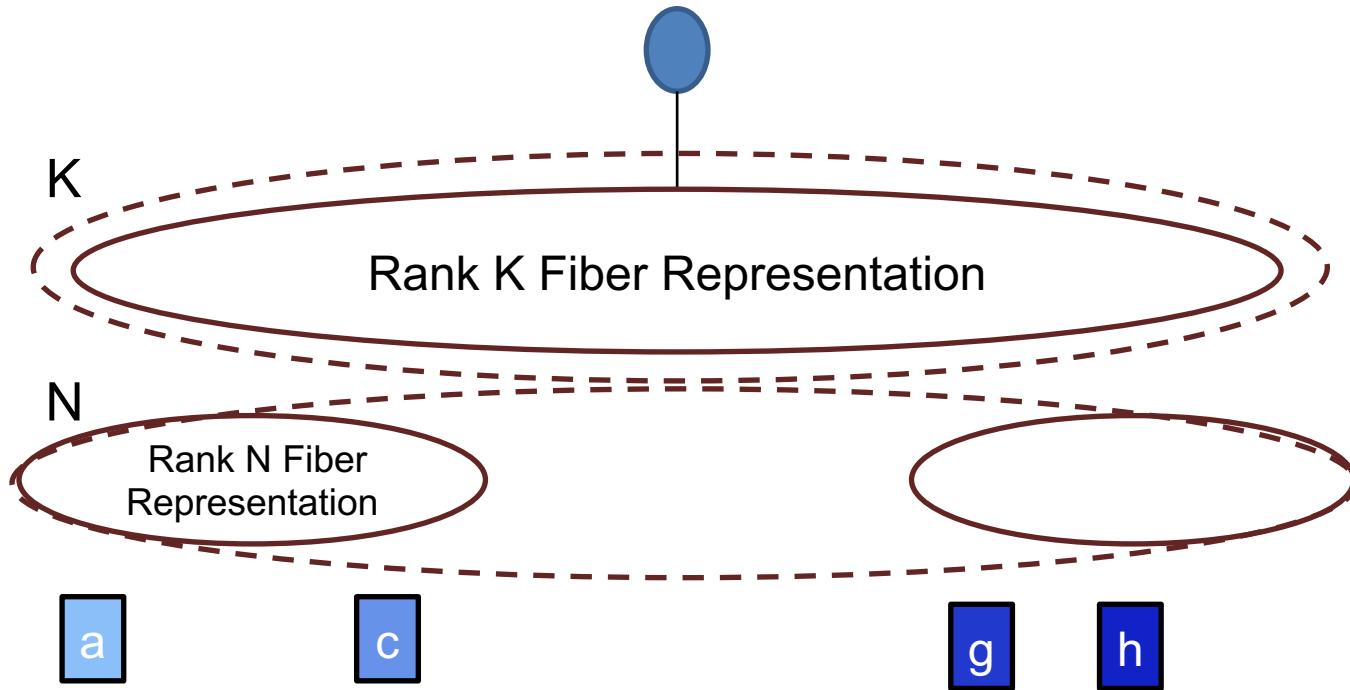
Information in a Fiber

- Each fiber has set of (coordinate, “payload”) tuples



Information in a Fiber

Method: `payload = fiber.lookup(coordinate)`



Fiber Representation Choices

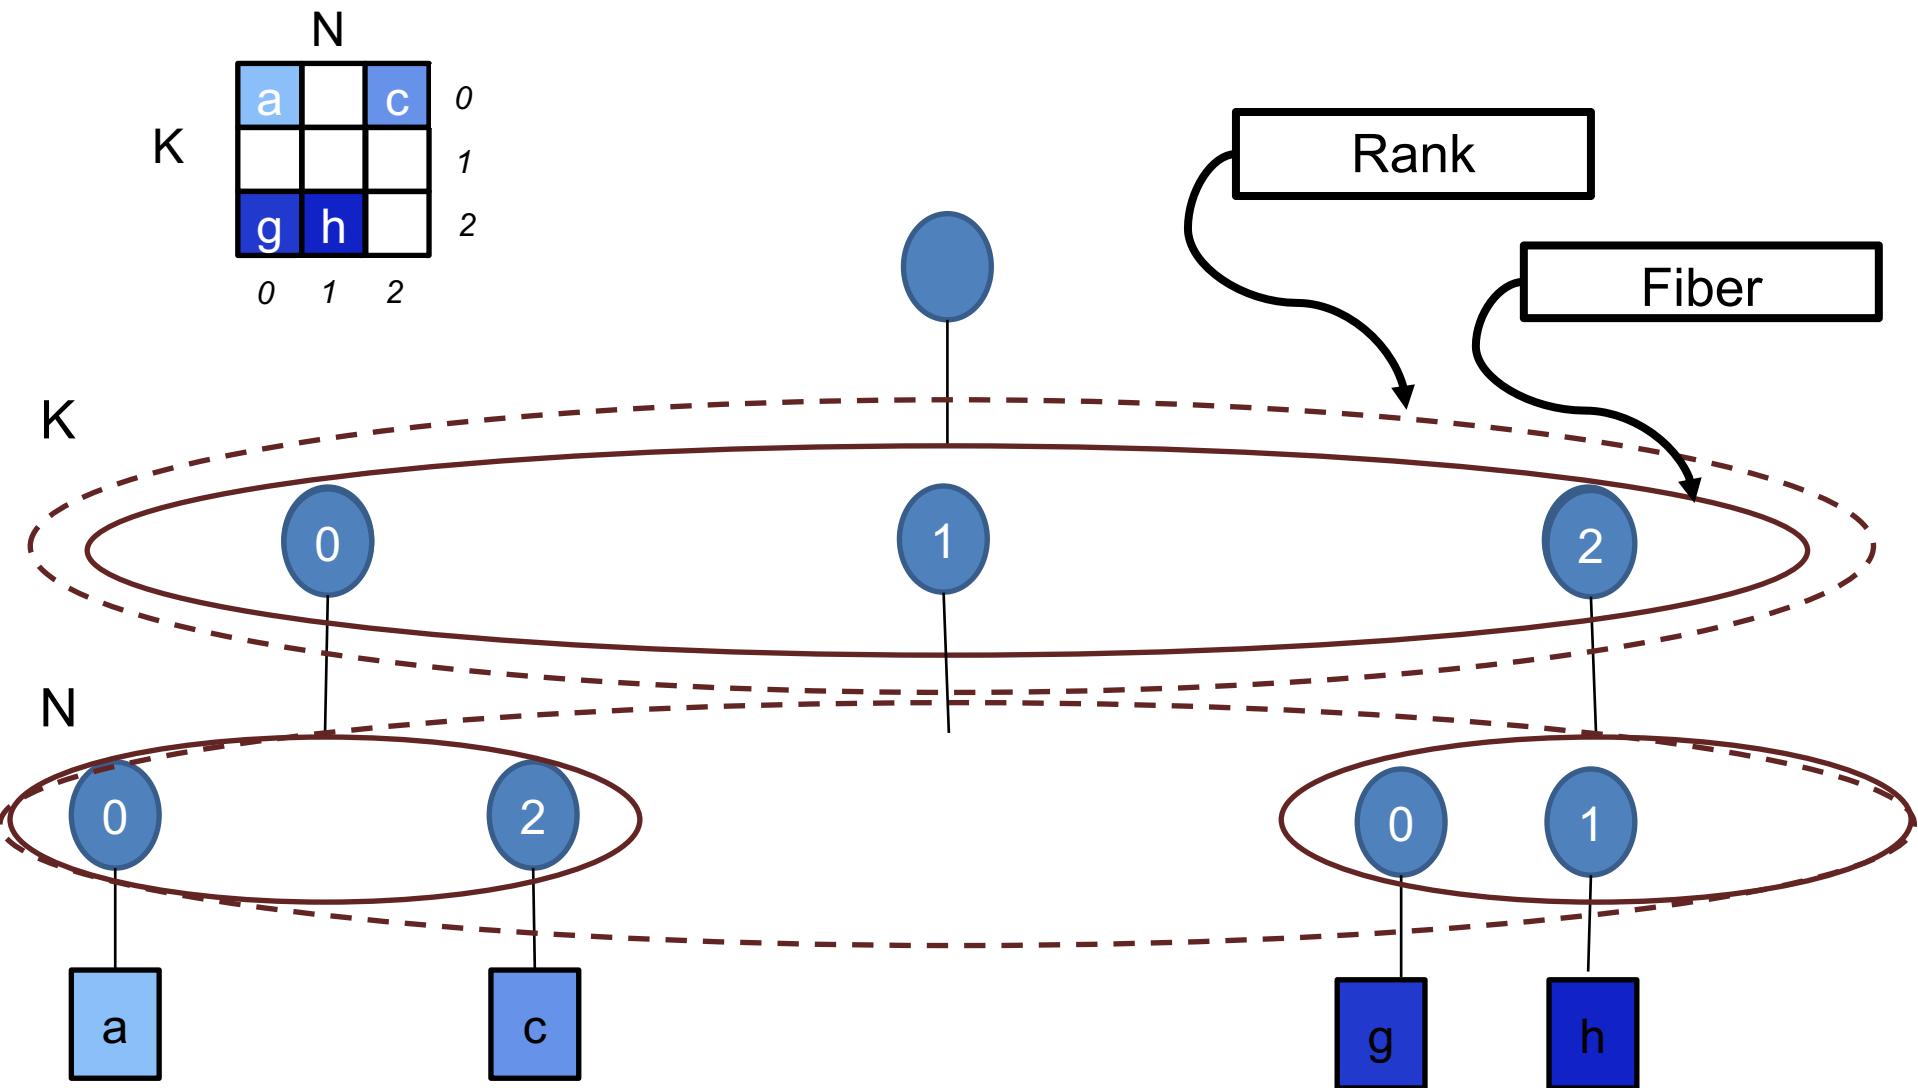
Each fiber has set of (coordinate, “payload”) tuples

- **Implicit Coordinates**
 - Uncompressed (no meta-data required)
 - Compressed – e.g., run length encoded
- **Explicit Coordinates**
 - E.g., coordinate list
- **Space efficiency of a representation depends on sparsity**
 - Compressed format can have **overhead** relative to uncompressed format for dense data

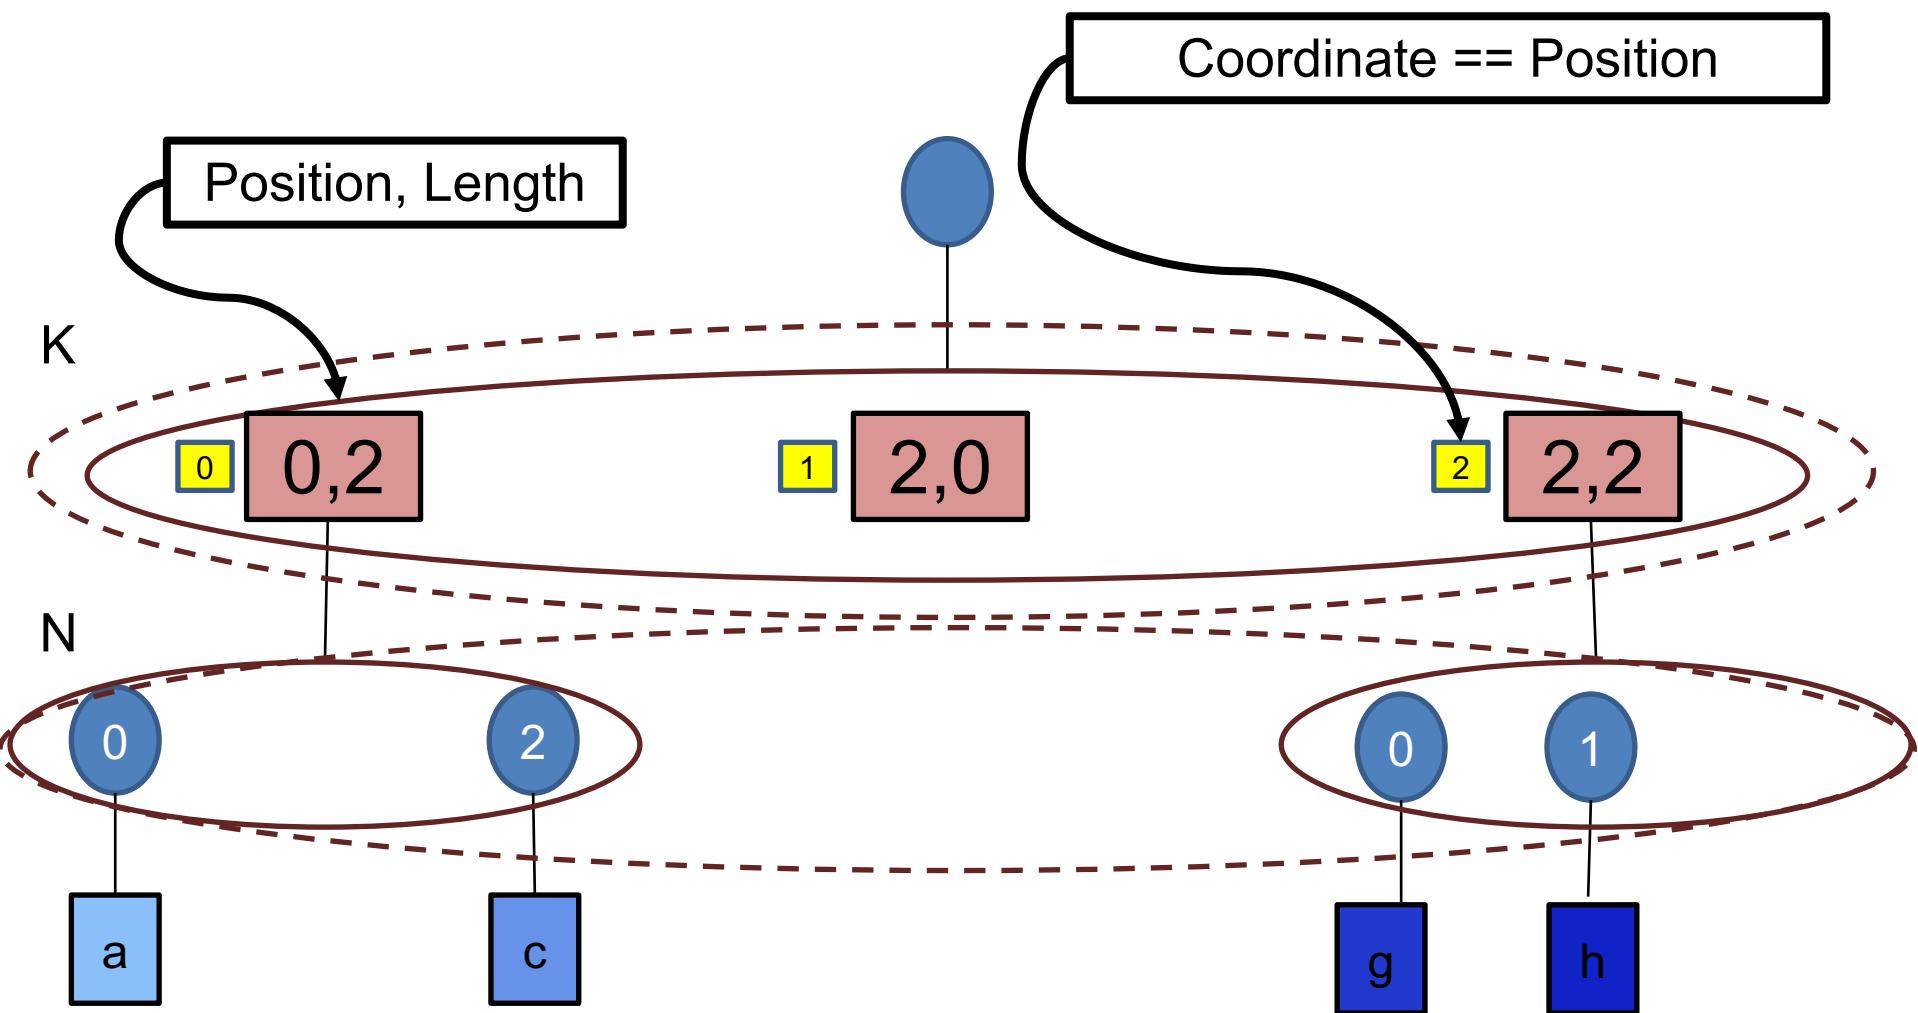
Compressed Implicit Coordinate Representations

- “Empty” coordinate compression via zero-run encoding
 - Run-length coding (RLE)
 - (**run-length of zeros**, **non-zero payload**)...
 - Significance map coding
 - (**flag to indicate if non-zero**, **non-zero payload**)...
- Payload encoding
 - Fixed length payload
 - Variable length payload
 - E.g., Huffman coding
- Efficiency of different traversal patterns through the tensor is affected by encoding, e.g., finding the payload for a particular coordinate...

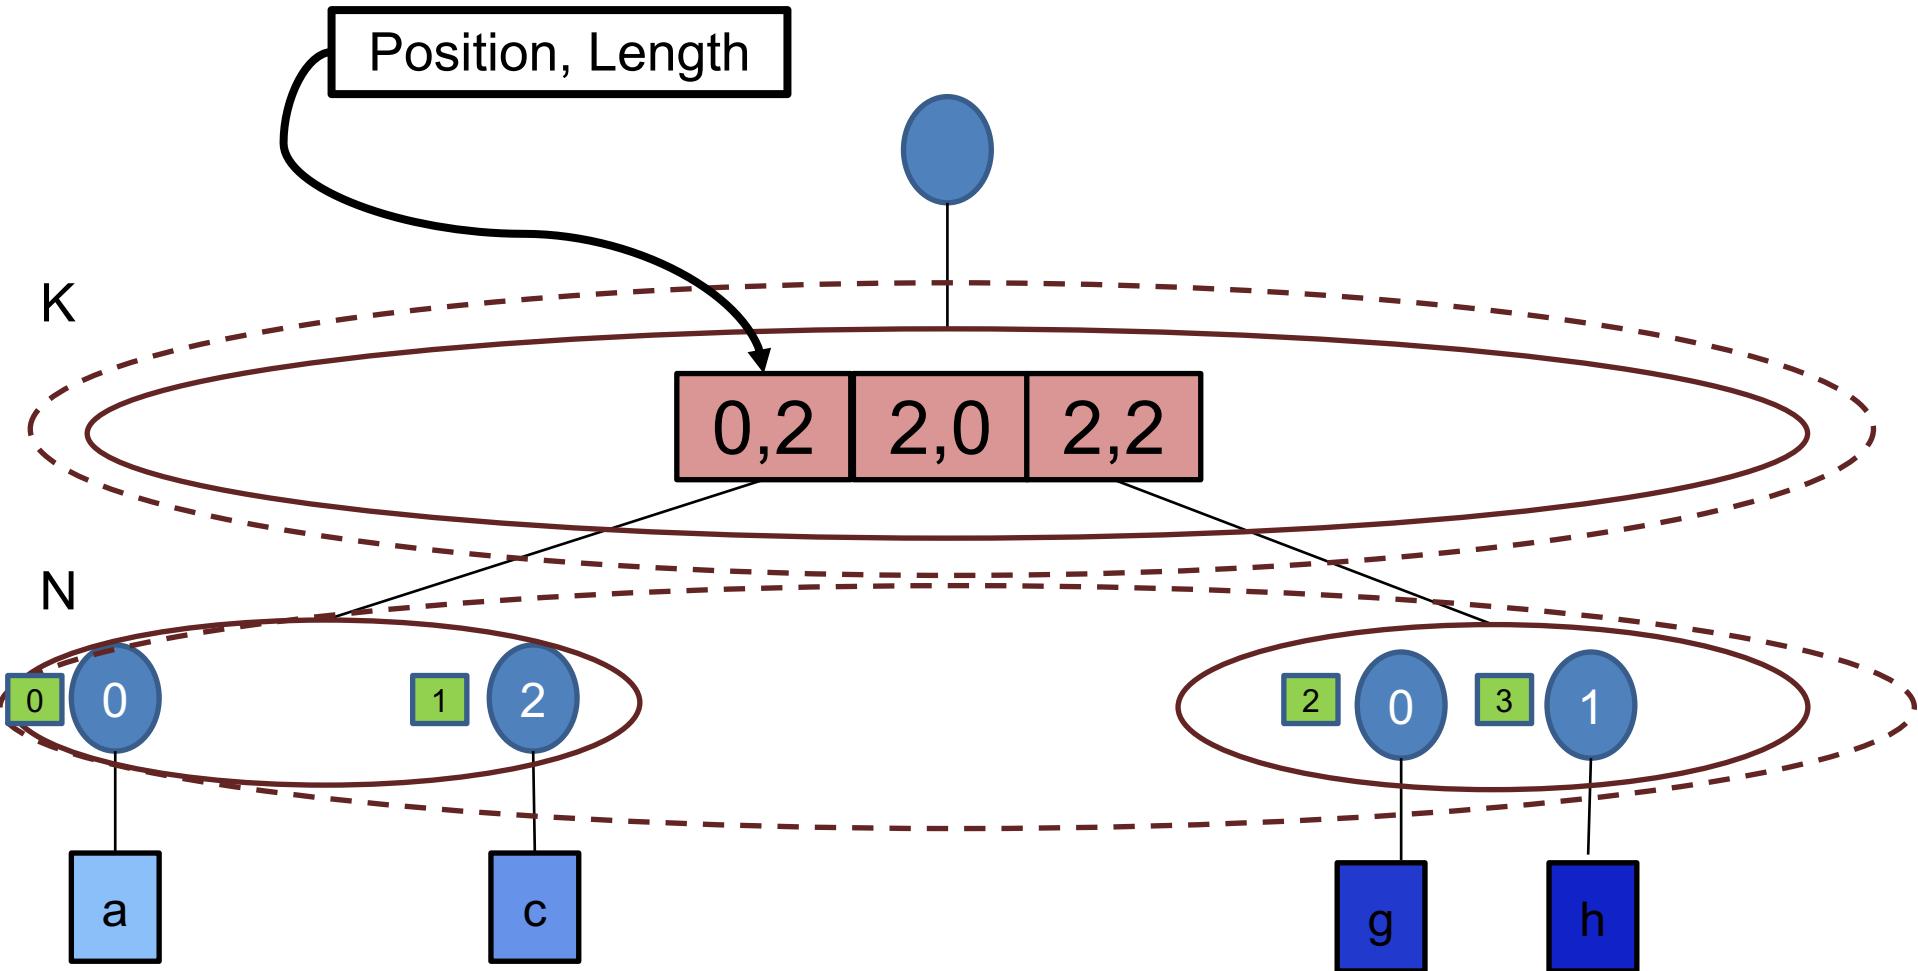
Uncompressed/Compressed Representation



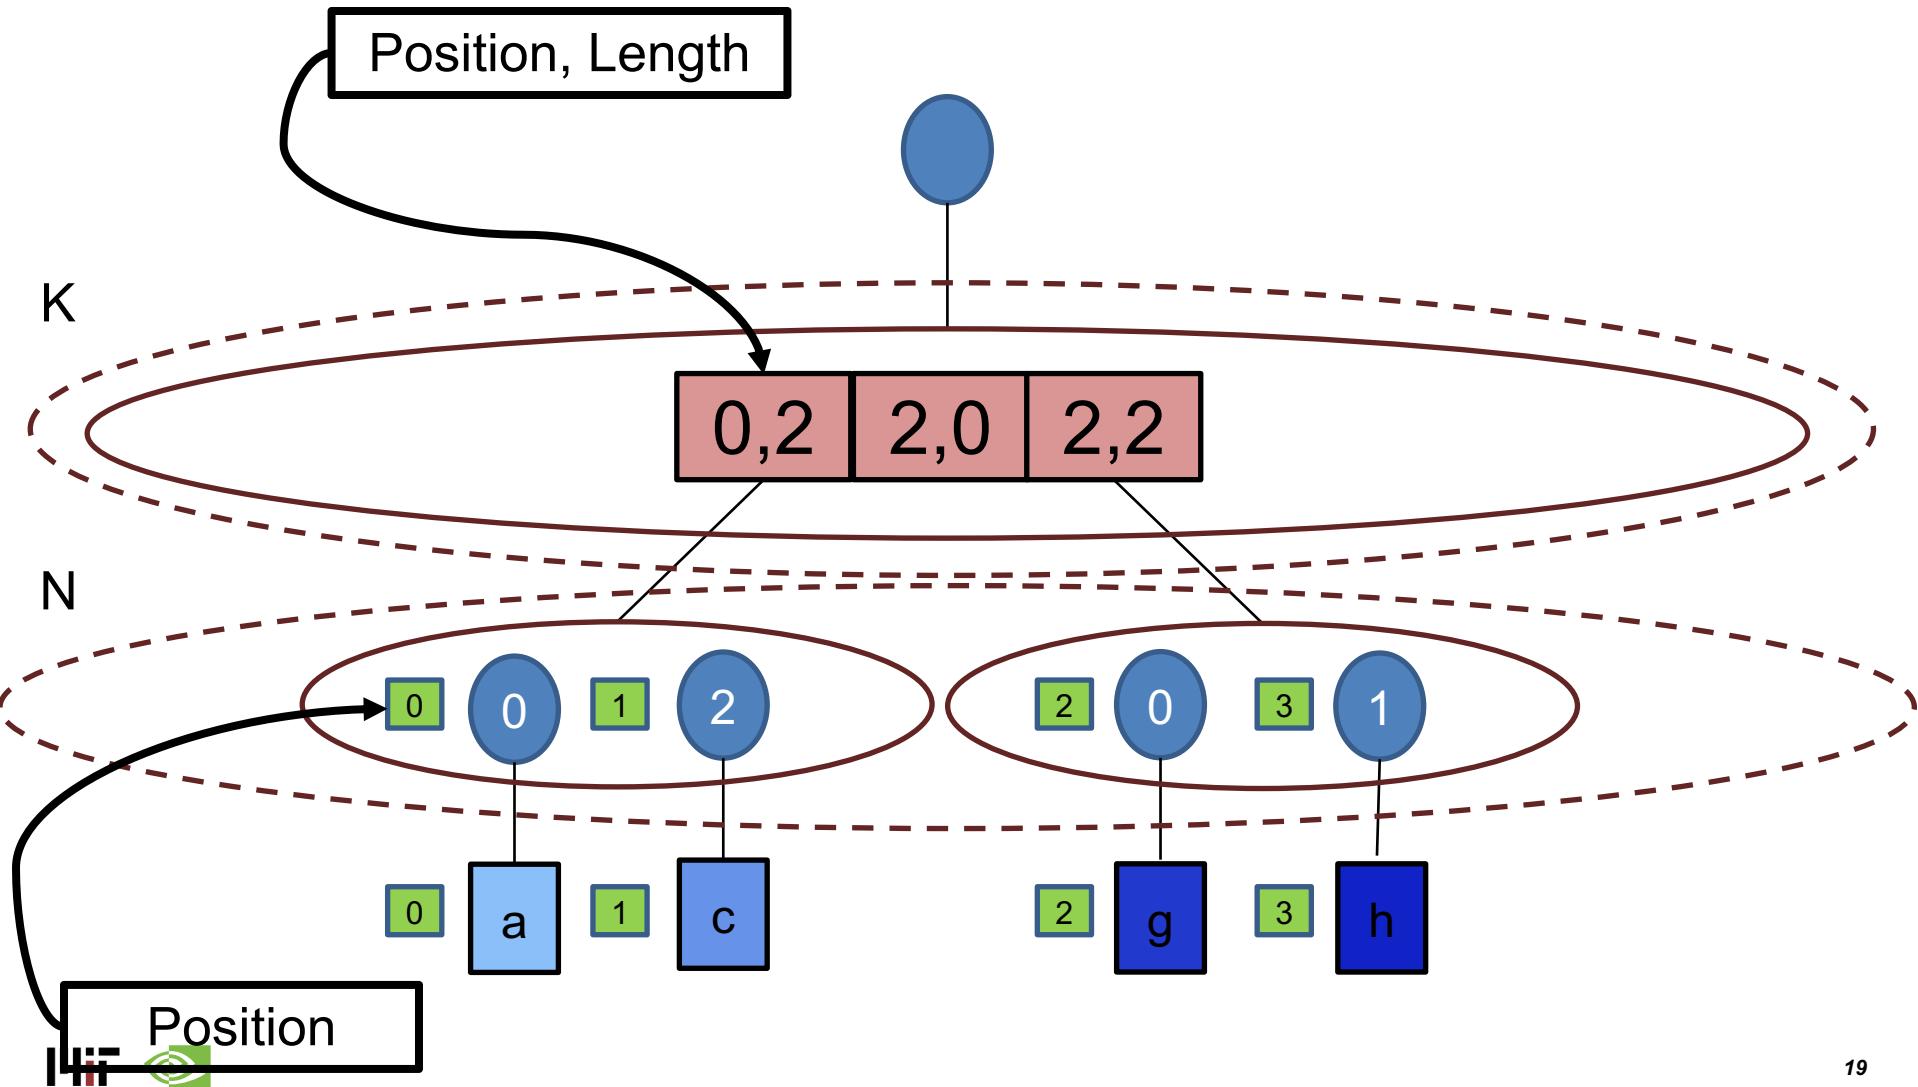
Uncompressed/Compressed Representation



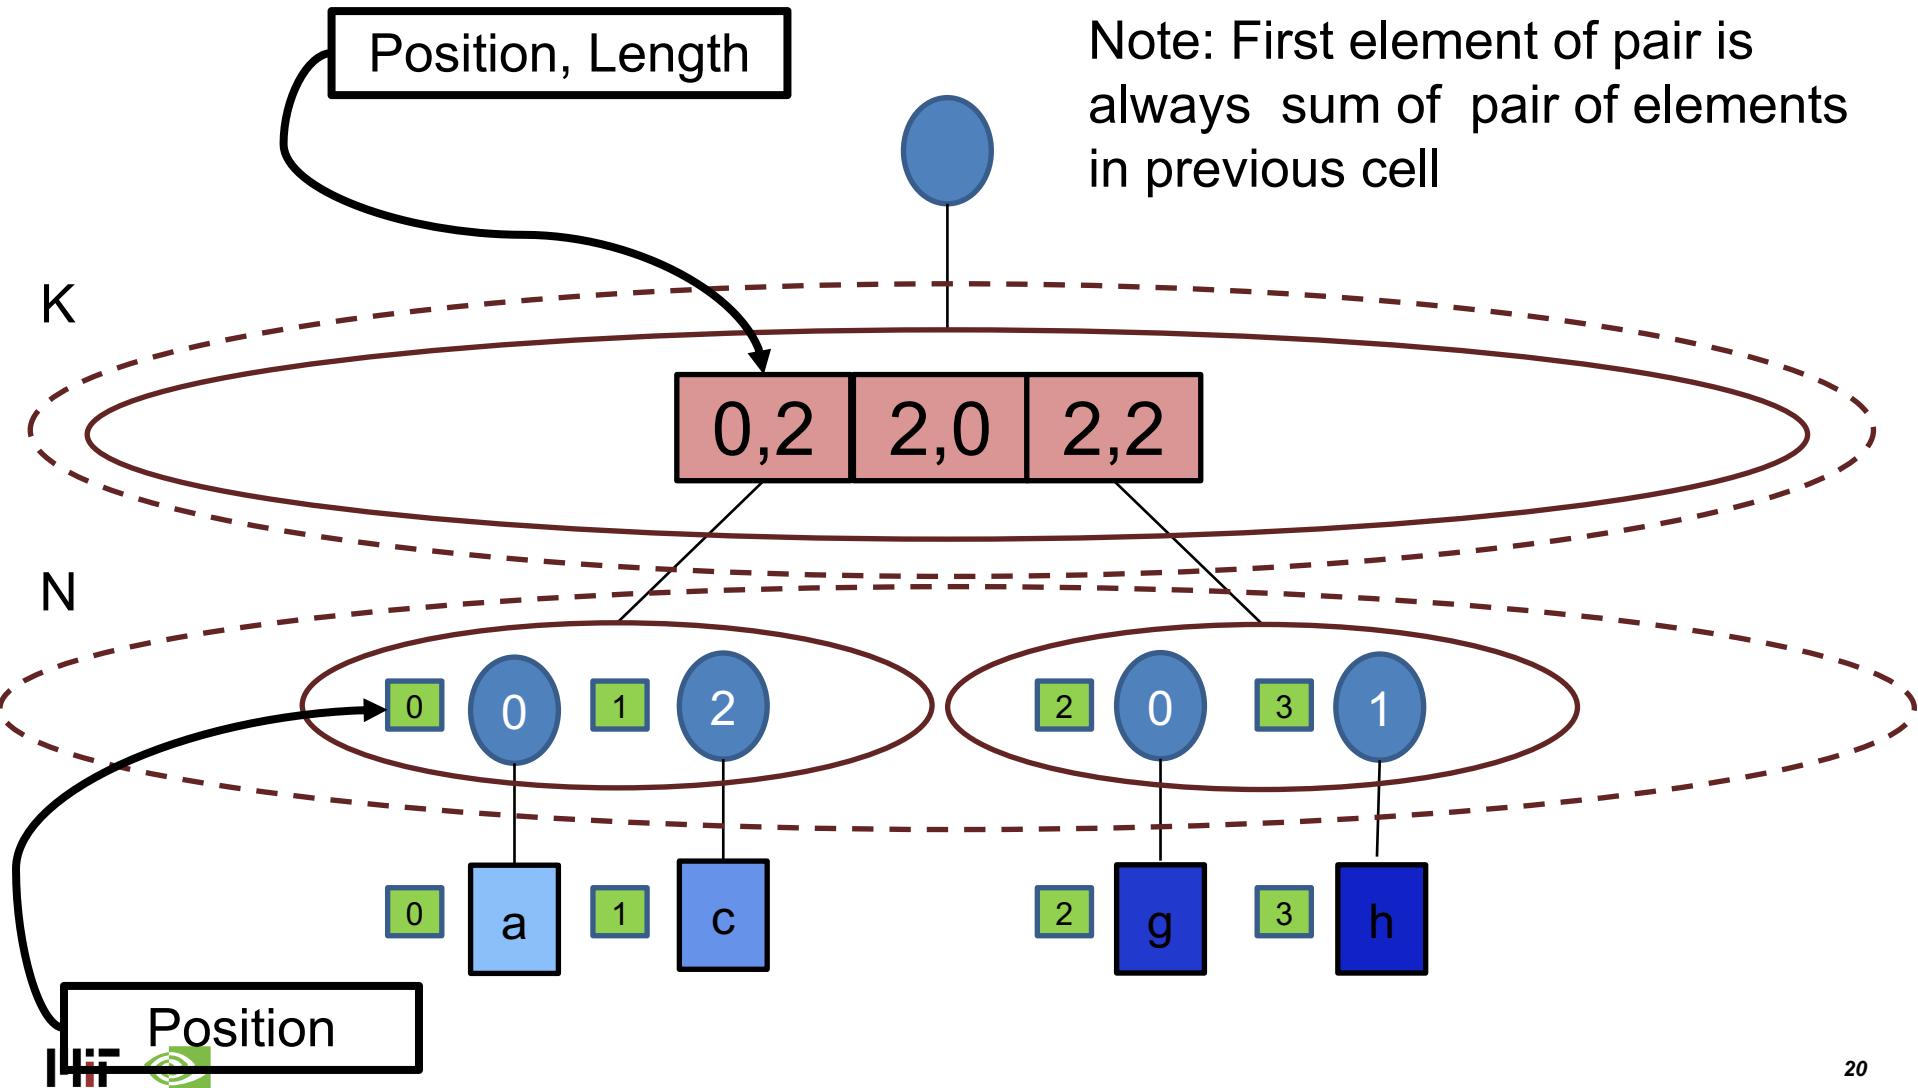
Uncompressed/Compressed Representation



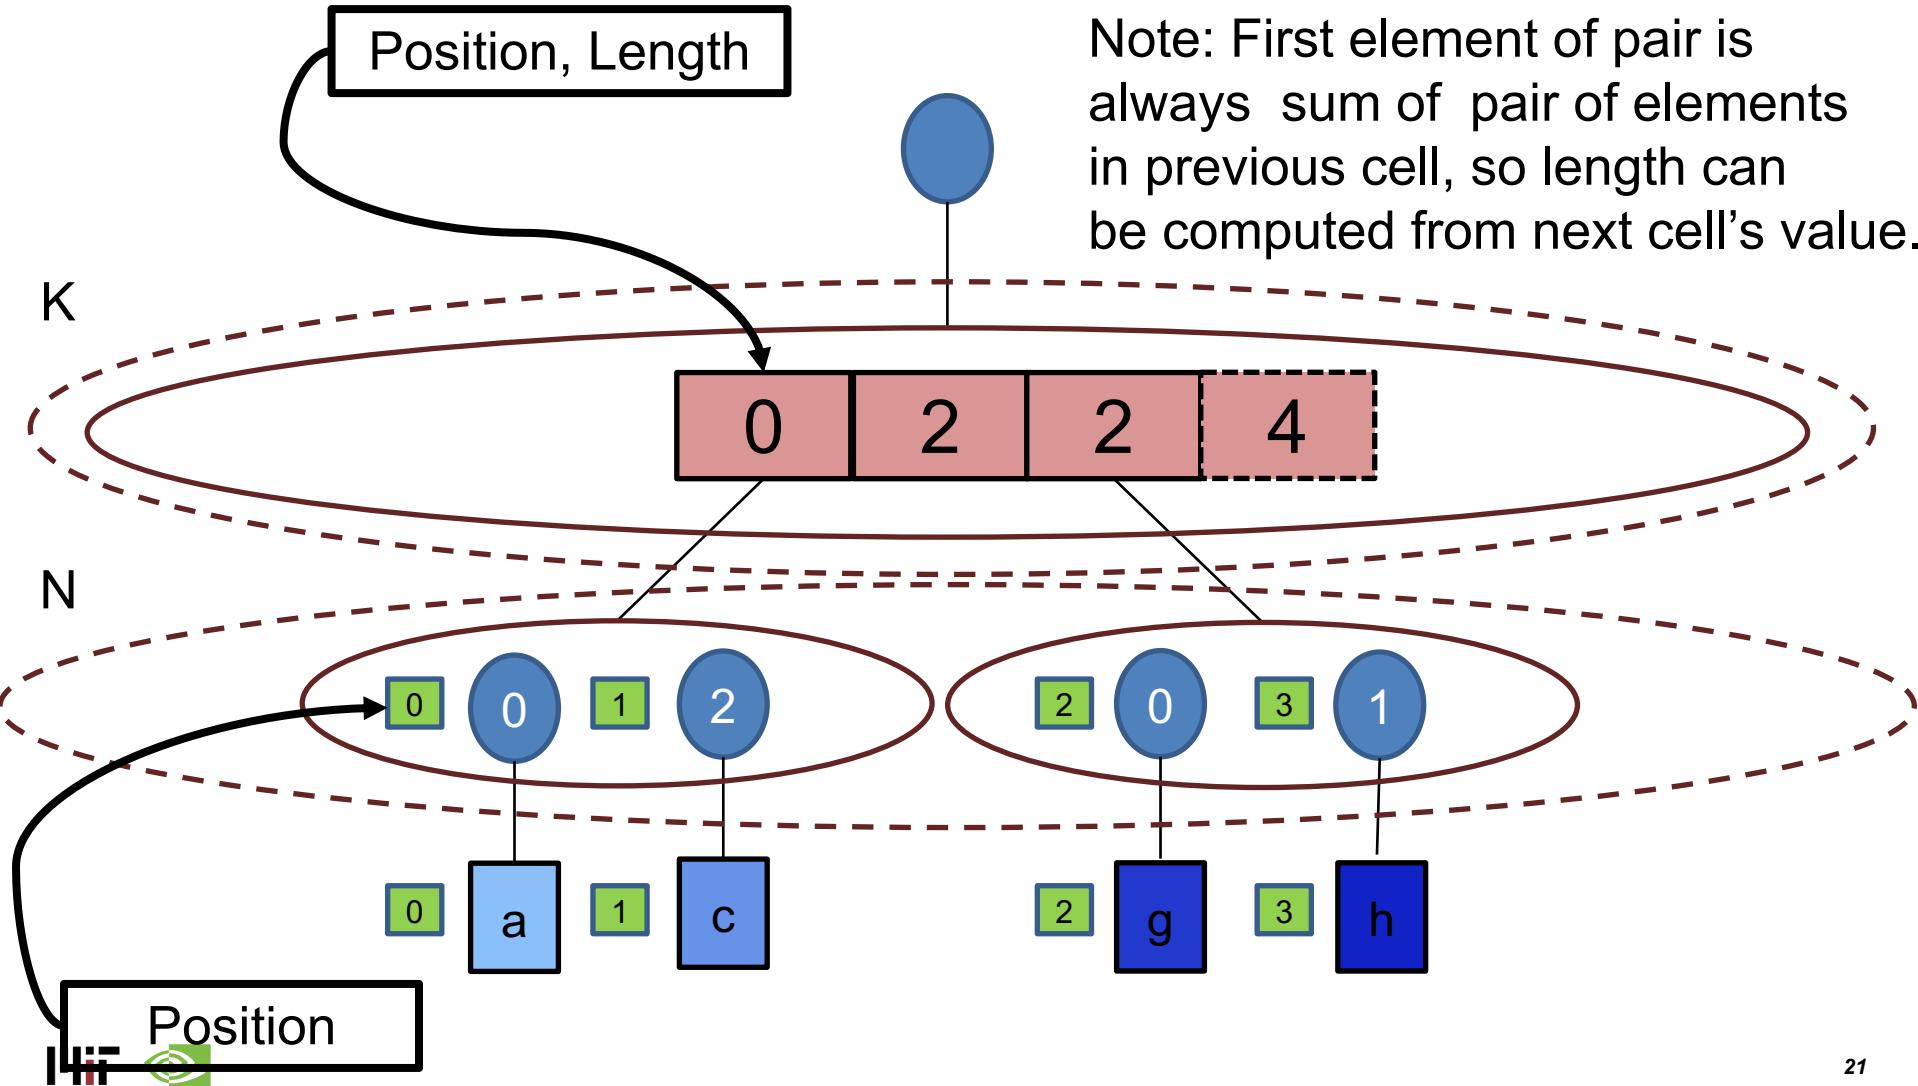
Uncompressed/Compressed Representation



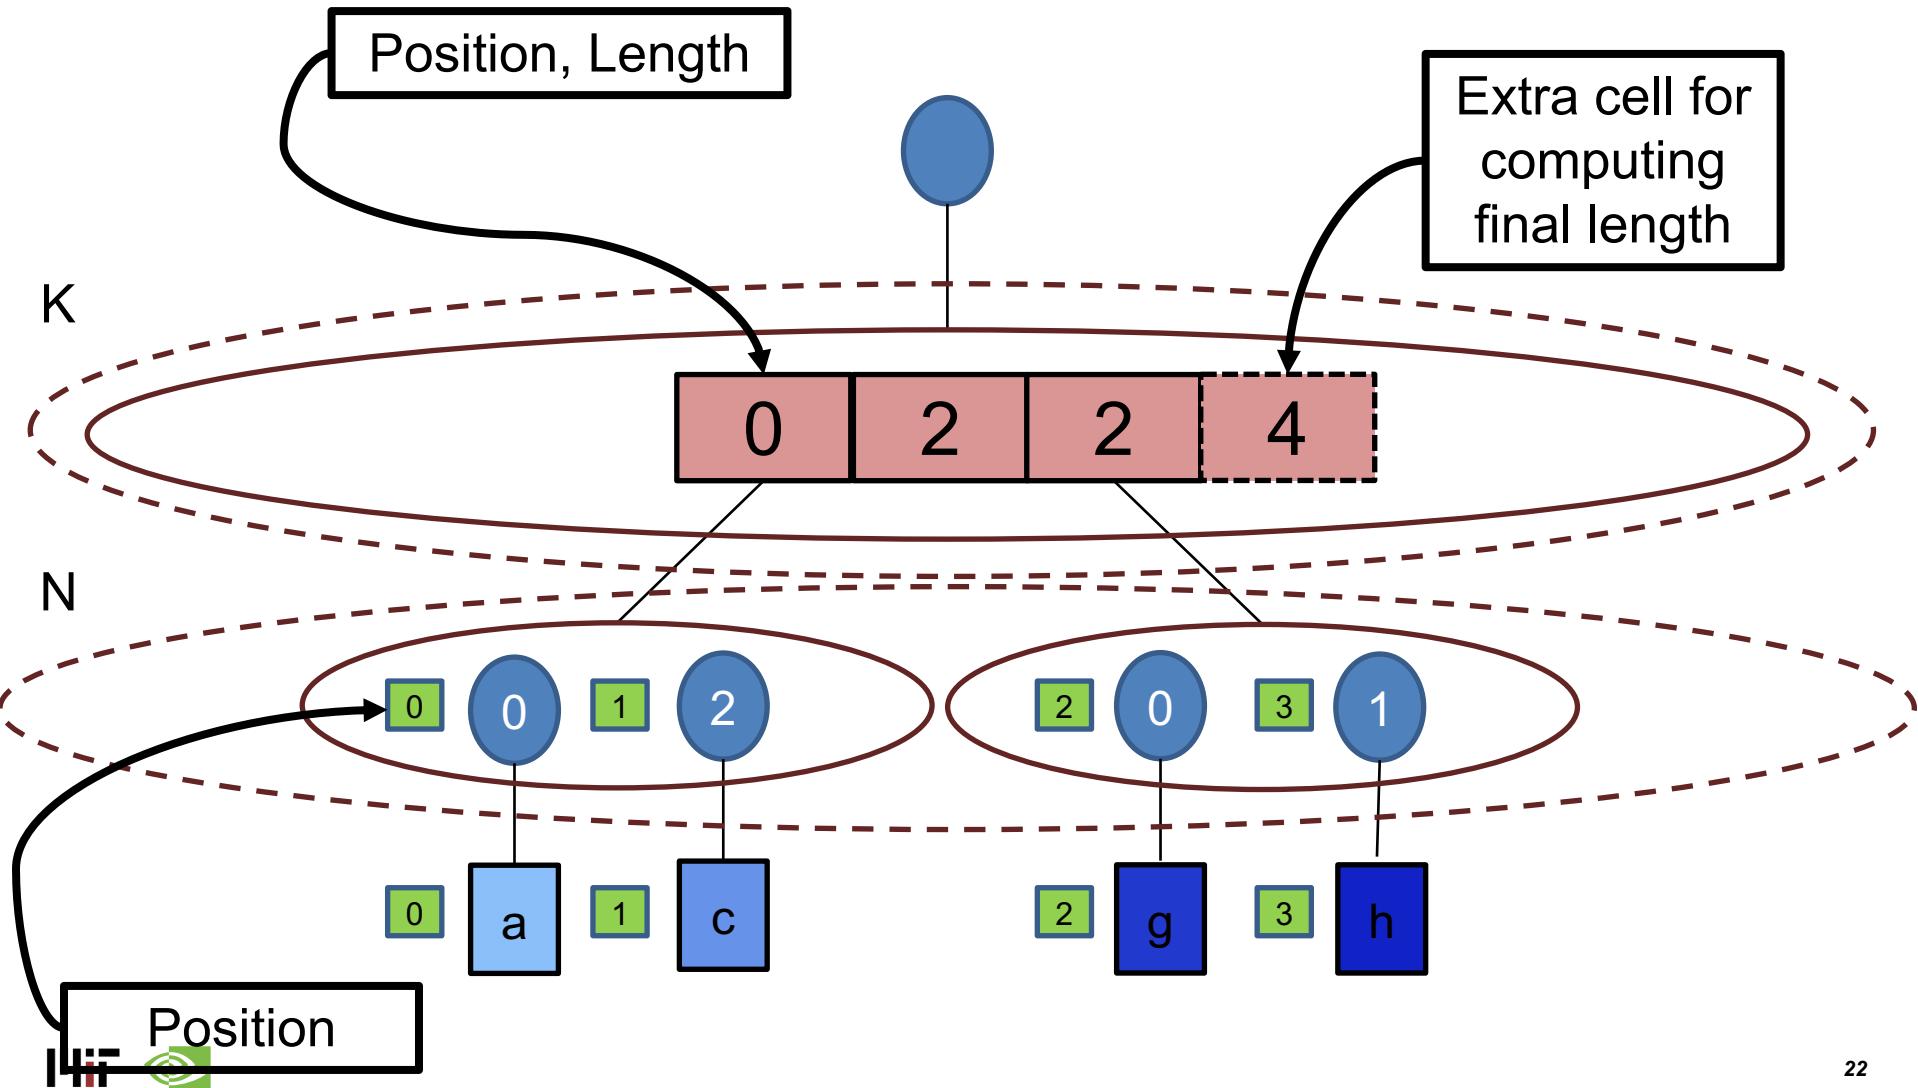
Uncompressed/Compressed Representation



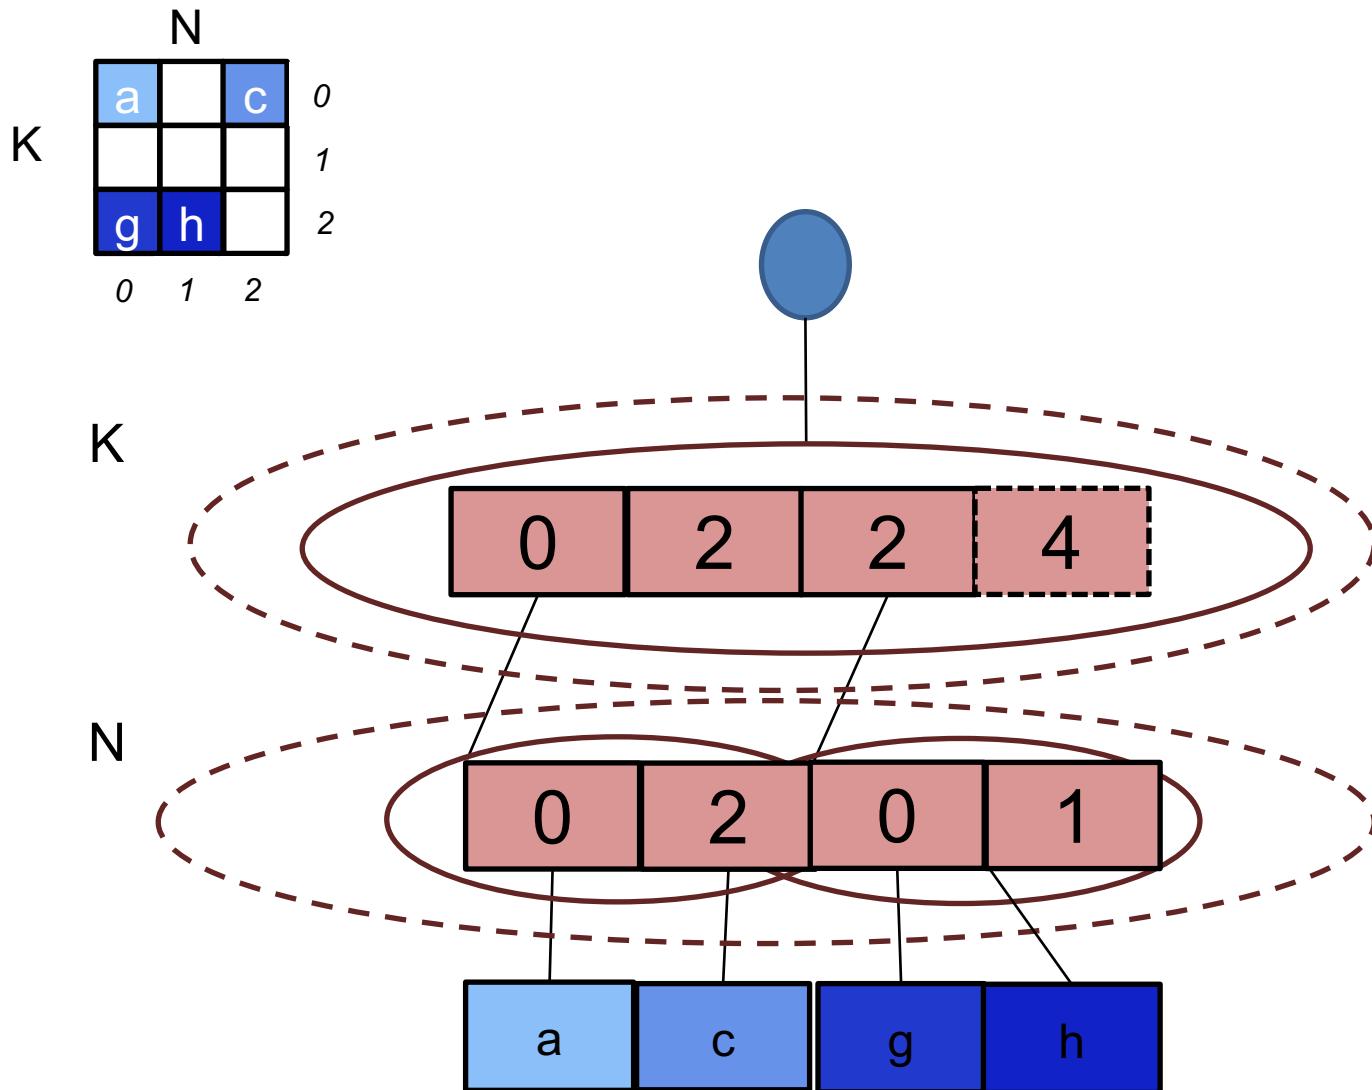
Uncompressed/Compressed Representation



Uncompressed/Compressed Representation

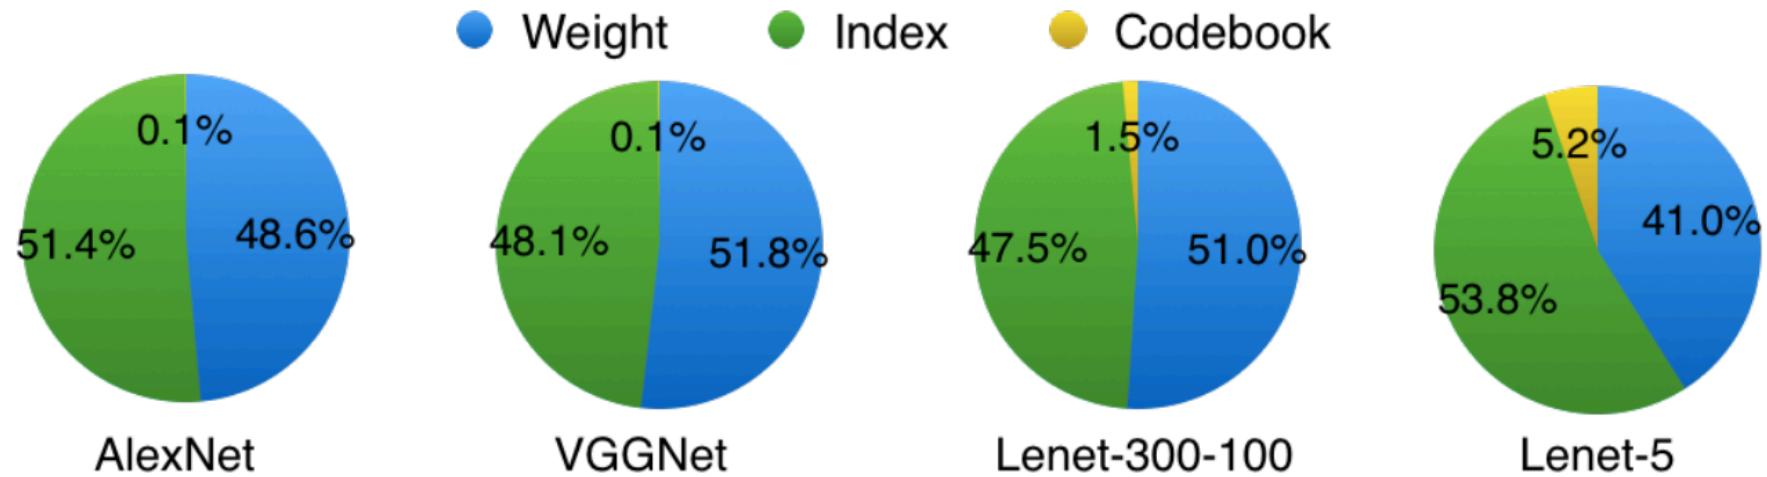


Compressed Sparse Row (CSR)



Compression Overhead

*Index (non-zero position info – e.g., **IA** and **JA** for CSR) accounts for approximately half of storage for fine grained pruning*



[Han et al., ICLR 2016]

Explicit Coordinate Representations

- Coordinate/Payload list
 - (**coordinate**, **non-zero payload**)...
- Hash table (per fiber)
 - (**coordinate** \rightarrow **payload**) mapping
- Hash table (per rank)
 - (**fiber_id**, **coordinate** \rightarrow **payload**) mapping
- Bit vector of non-zero coordinates
 - Uncompressed payload

Per Rank Tensor Representations

- Uncompressed [U]
 -
- Run-length Encoded [R]
 -
- Coordinate/Payload List [C]
 -
- Hash Table (per rank) [H_r]
- Hash Table (per fiber) [H_f]
 -

Inspired by collaboration with Kjolstad
in [Kjolstad, TACO]

Payload Needed to Find Fiber in a Rank

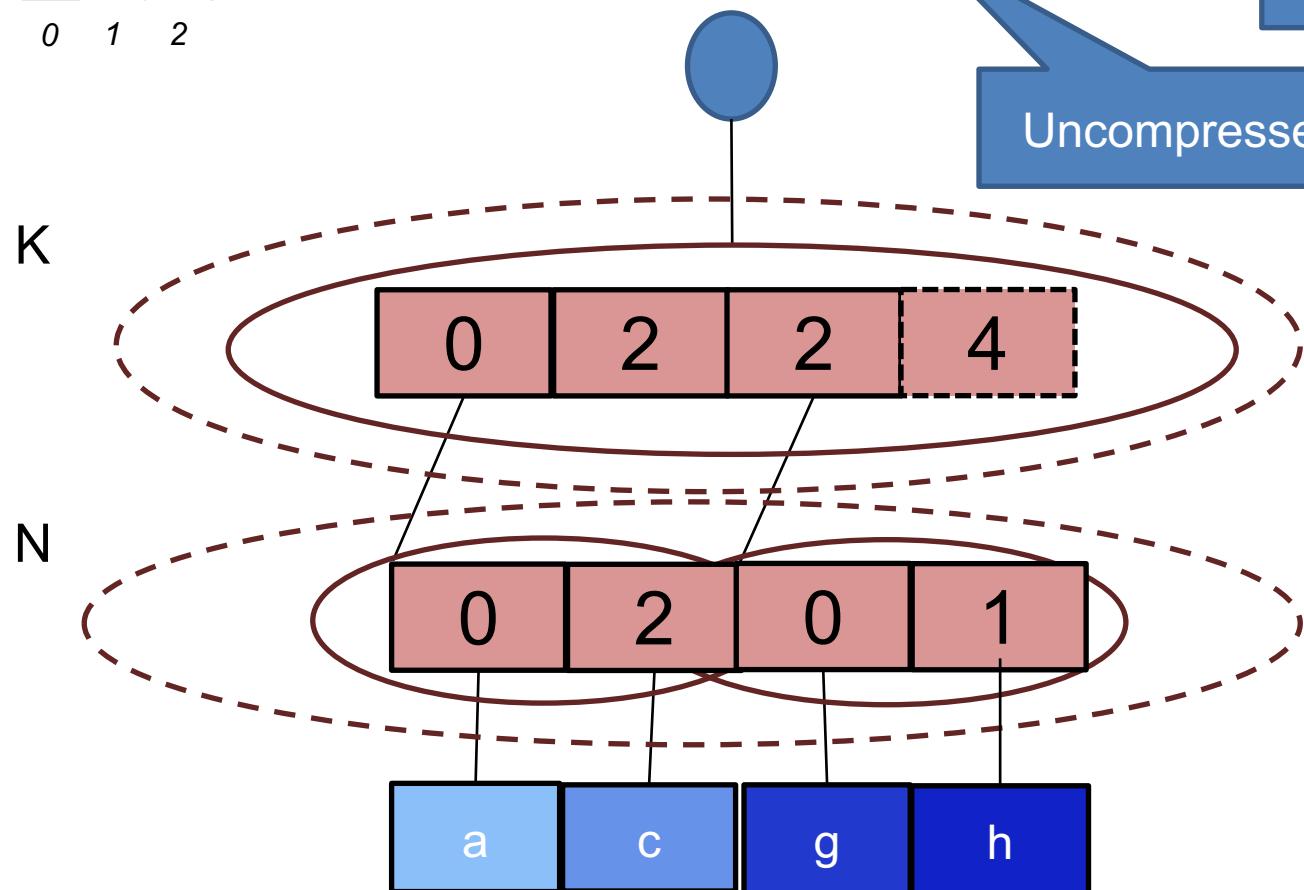
- Uncompressed [U]
 - Payload*: None
- Run-length Encoded [R]
 - Payload*: Pointer to fiber data structure
- Coordinate/Payload List [C]
 - Payload*: Pointer to fiber data structure
- Hash Table (per rank) [H_r]
 - Payload*: fiber_id
- Hash Table (per fiber) [H_f]
 - Payload*: Pointer to fiber data structure

*Payload needed in preceding rank to perform “lookup():”

Notation for CSR

	N	
a		c
g	h	
0	1	2

CSR: Tensor-UC/KN

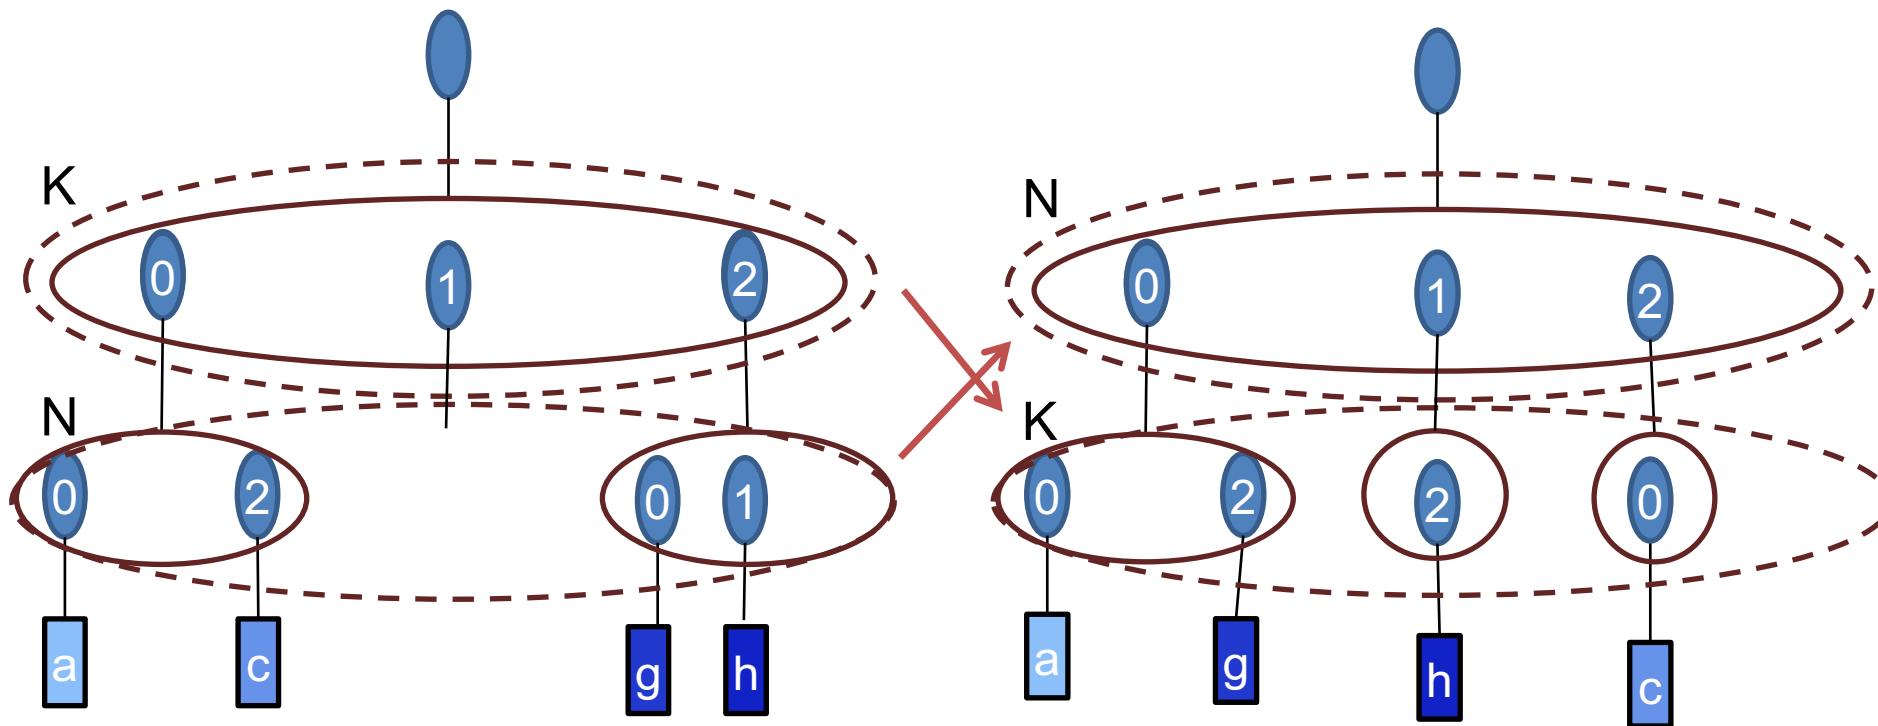


Representation of Order of Ranks

Tensor-UC/KN \rightarrow CSR

		N	
		a	c
		g	h
0			
1			
2			

Tensor-UC/NK \rightarrow CSC



Traversal Efficiency

Efficiency of different traversal patterns through the tensor is affected by encoding, e.g., finding the payload for a particular coordinate...

- Operations:
 - `payload = Tensor.Locate(coordinate... | point)`
 - `(coordinate, payload) = Tensor.Next(rank_traversal_order)`

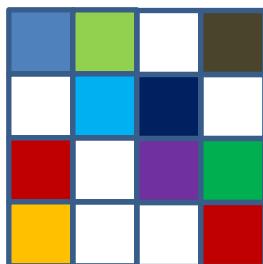
`Tensor.next()` is a very common operation and its efficiency is highly dependent on representation, both order of ranks and representation of each rank....

Concordant traversal orders

CSR and CSC each has a natural (or “concordant”*) traversal order

Processing
Order

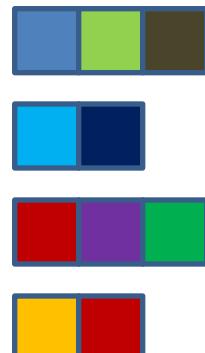
Original
Matrix



Row-major order

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}$$

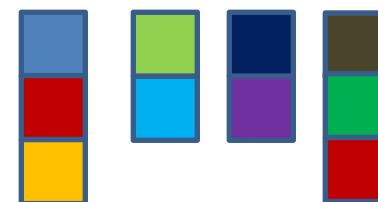
Compressed
Sparse Row
(CSR)



Column-major order

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}$$

Compressed
Sparse Column
(CSC)



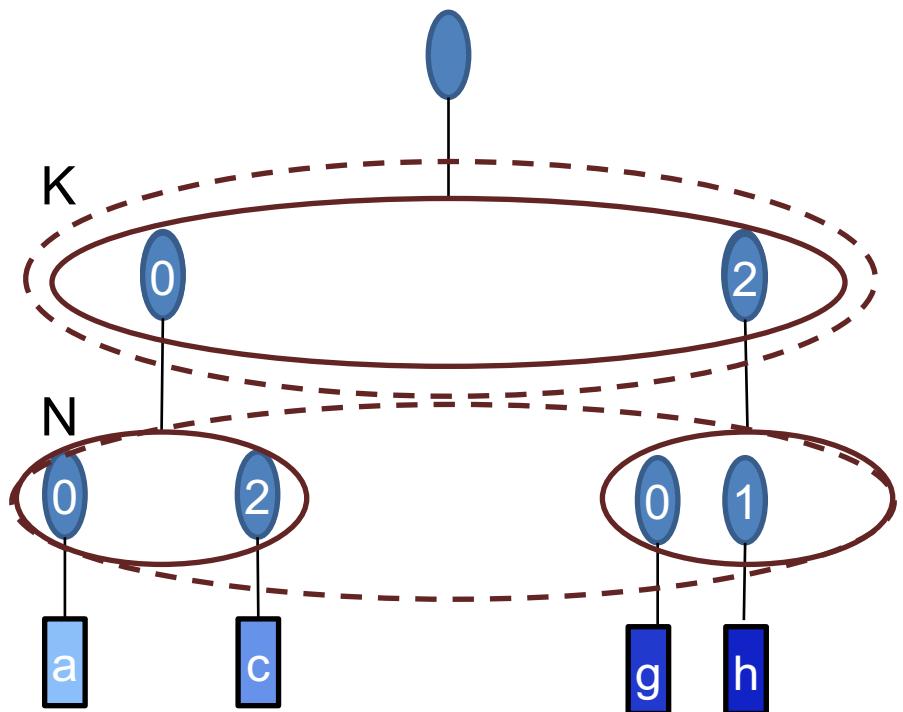
* Term from Michael Pellauer

Example Traversal Efficiency

- Locate efficiency:
 - Uncompressed – direct reference - O(1)
 - Run length encoded – linear search – O(n)
 - Hash table – multiple references and compute – O(1)
 - Coordinate/Payload list – binary search – O(log n)
- Next efficiency (concordant traversal)
 - Uncompressed – sequential reference, good spatial locality - O(1)
 - Run length encoded – sequential reference – O(1)
 - Coordinate/Payload list - same as uncompressed
- Next efficiency (discordant traversal)
 - Essentially as good (or bad) as locate....

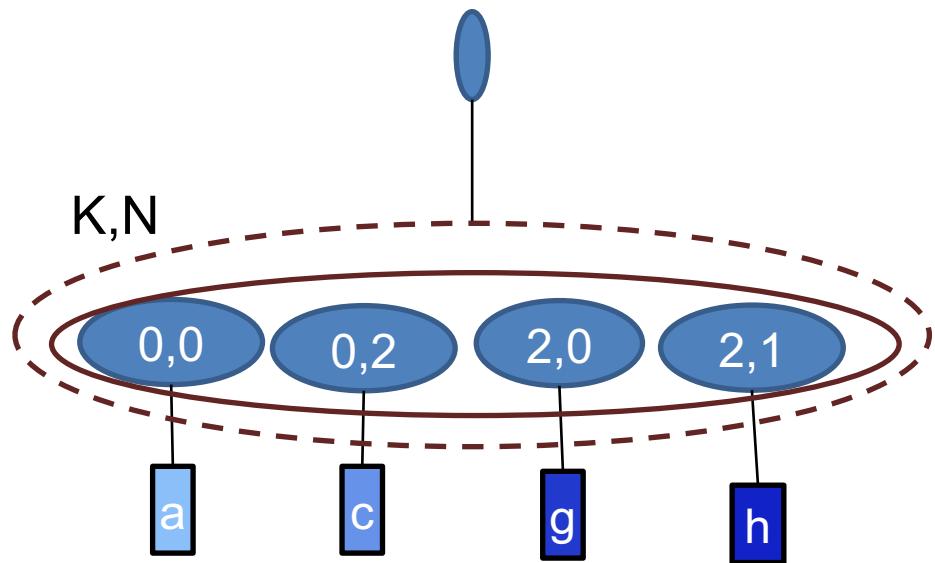
Merging Ranks

Tensor-CC/KN



		N	
		K	
0	1	a	c
1	2		
2		g	h

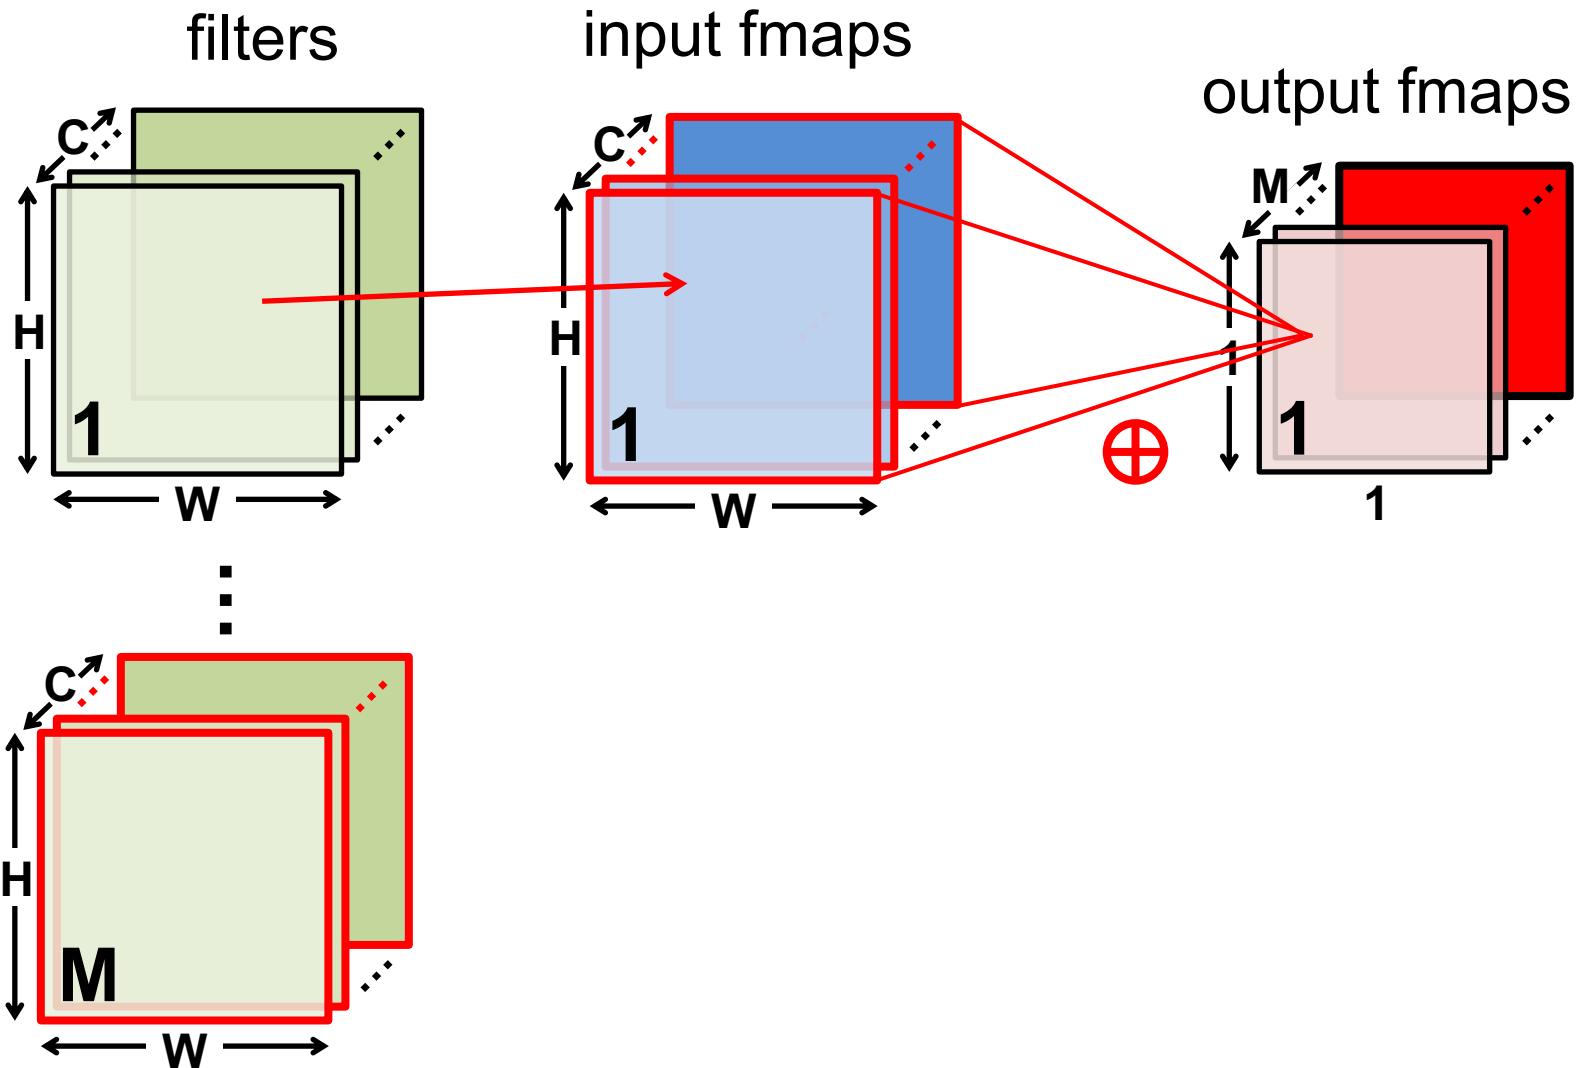
Tensor- $(C^2)/(KN)$



Merging Ranks

- For efficiency one can form new representations where the data structure for two or more ranks are combined:
- Examples:
 - Tensor-(C²)
List of (coordinate tuple,payload) - COO
 - Tensor-(H²)
 - Hash table with coordinate tuple as key
 - Tensor-(U²)
 - Flattened array
 - Coordinates can be recovered with modulo arithmetic on “position”
 - Tensor-(R²)
 - Flattened run-length encoded sequence

Fully-Connected (FC) Layer



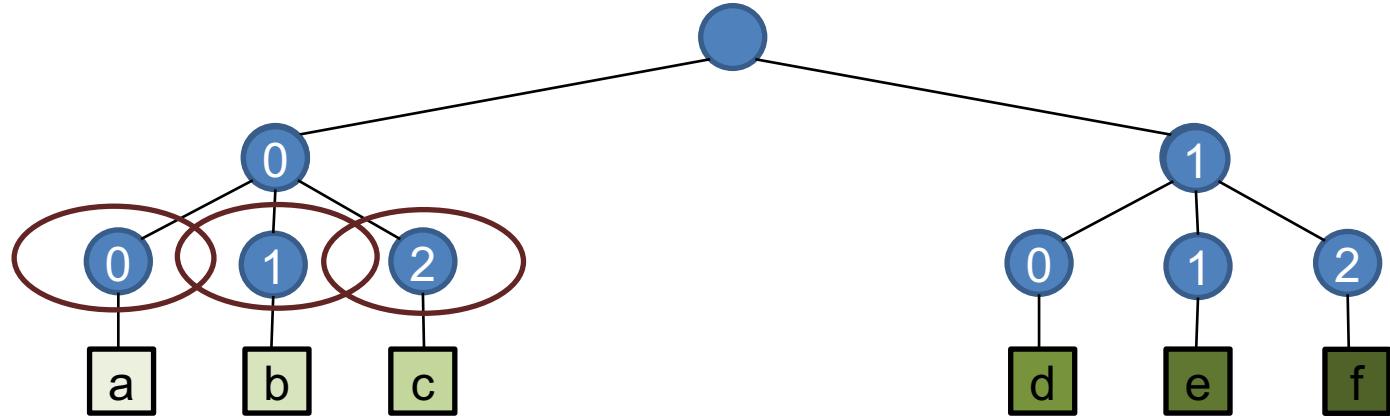
Output-Stationary Operations

Filter

a	b	c
d	e	f

M

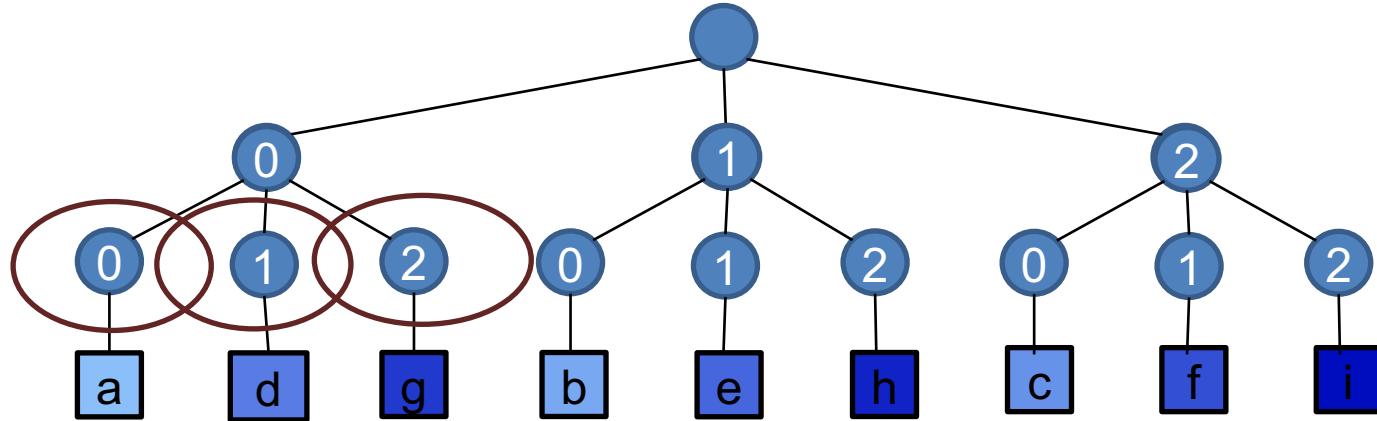
K



K

a	b	c
d	e	f
g	h	i

N



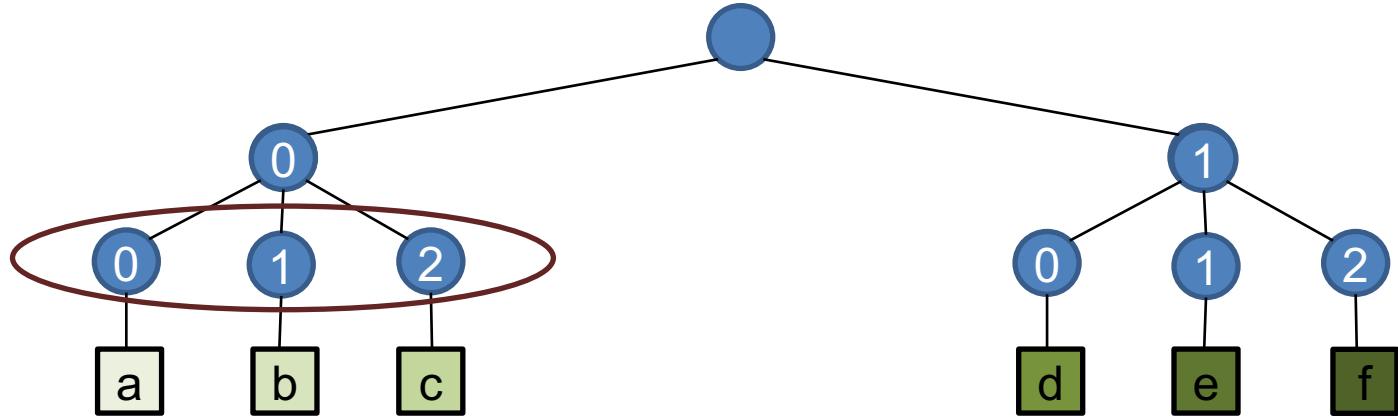
Output-Stationary Lists

Filter

a	b	c
d	e	f

M

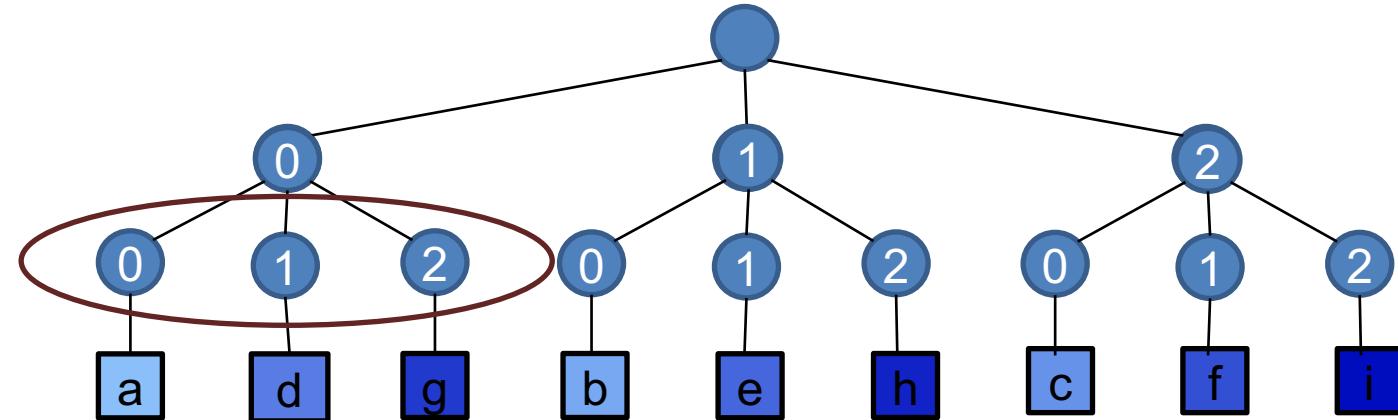
K



K

a	b	c
d	e	f
g	h	i

N



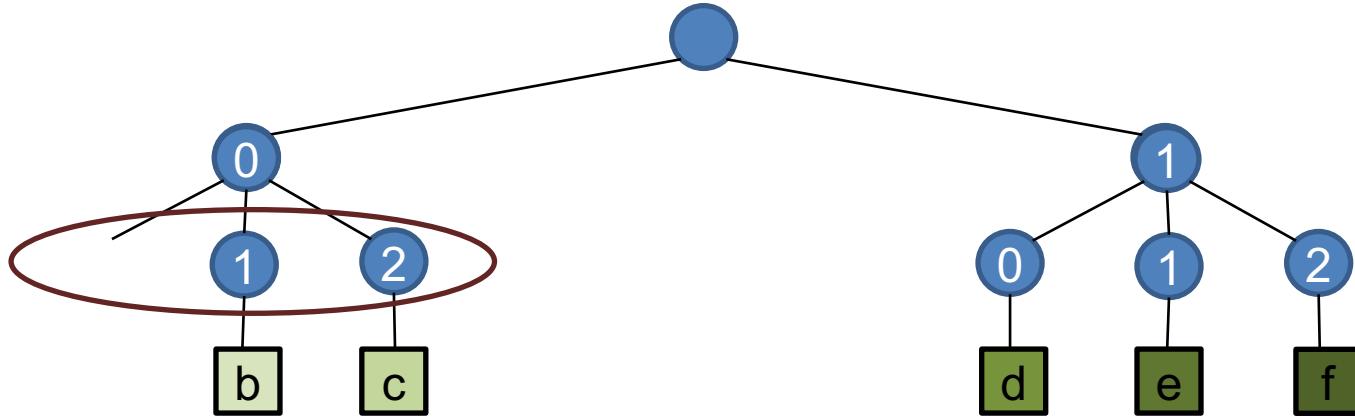
Output-Stationary Intersection Lists

Filter

	b	c
d	e	f

M

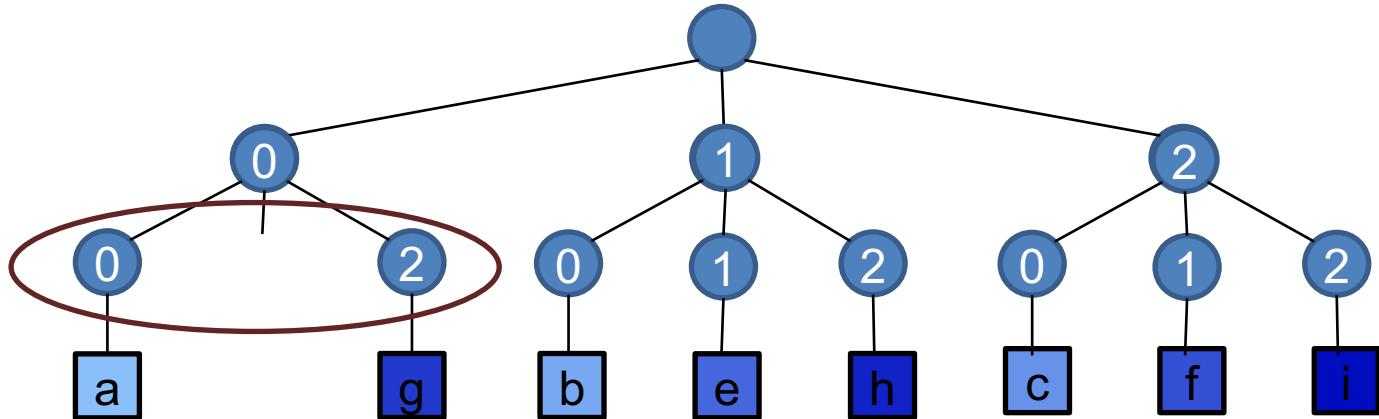
K



K

a	b	c
e	f	
g	h	i

N

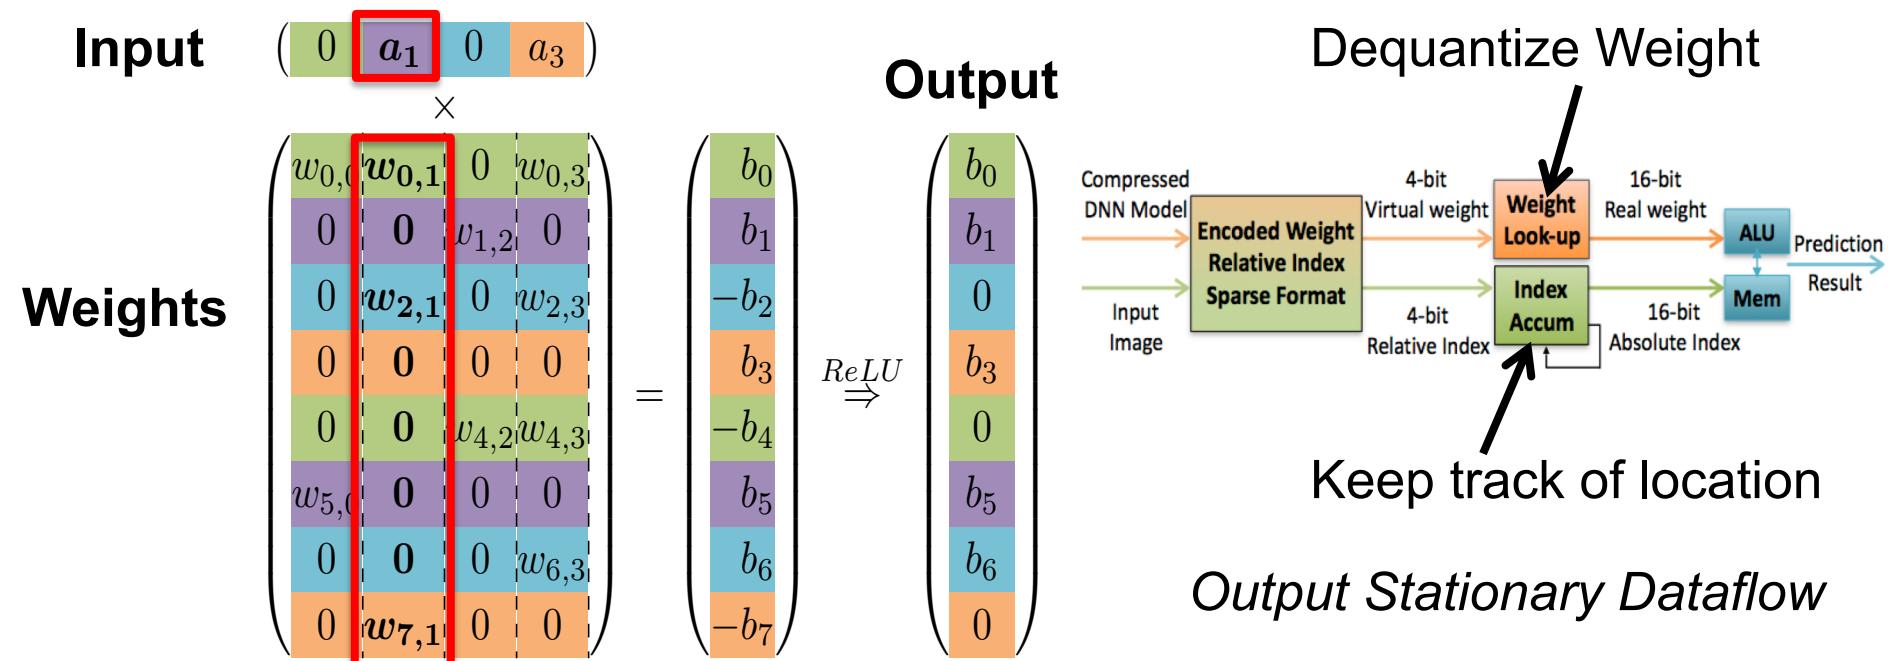


2-2 is the only “effectual” computation

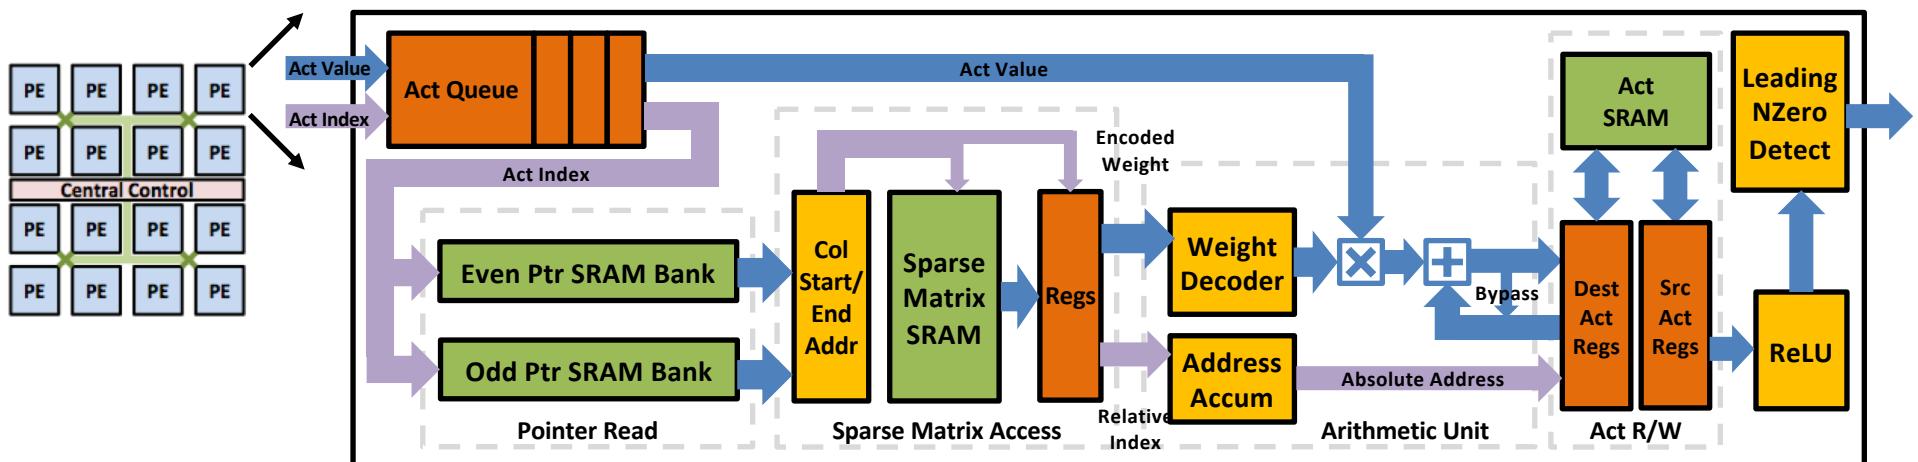
EIE: A Sparse Linear Algebra Engine

- Process Fully Connected Layers (after Deep Compression)
- Store weights column-wise in Run Length format (i.e., CSC format)
- Read relative column when input is non-zero

Supports Fully Connected Layers Only



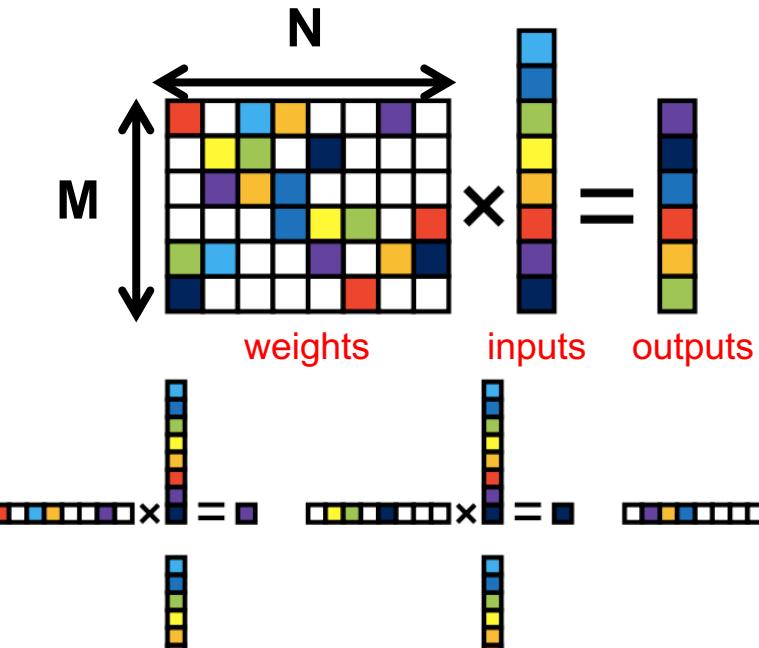
PE Architecture



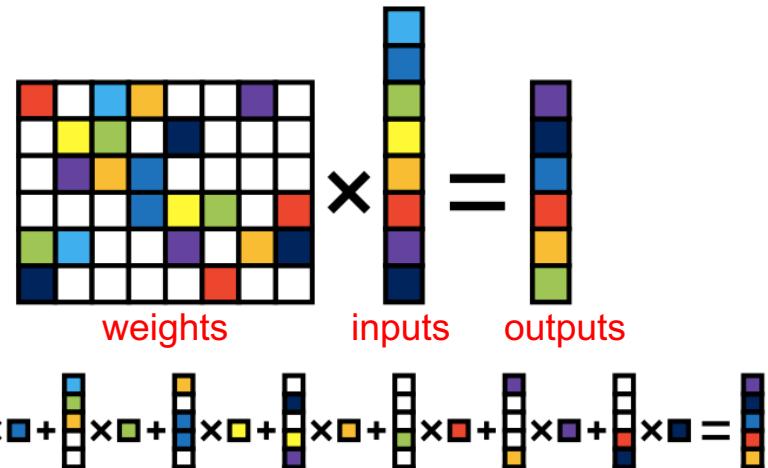
Impact of Representation on Dataflow

From SpMxV research

Compressed Sparse Row (CSR)



Compressed Sparse Column (CSC)



Output stationary

Input stationary

CSC reduces memory bandwidth over CSR
(when not $M \gg N$)

For DNN, $M = \# \text{ of filters}$,
 $N = \# \text{ of weights per filter}$

Sparse Accelerators

1-D Output-Stationary Convolution



```

int i[W];      # Input activations
int w[R];      # Filter weights
int o[E];      # Output activations

for (e = 0; e < E; e++) {
    for (r = 0; r < R; r++) {
        o[e] += i[e+r]*w[r];
    }
}

```

What opportunity(ies) exist if some of the values are zero?

Can avoid reading operands, doing multiply and updating output

[†] Assuming: ‘valid’ style convolution

1-D Output-Stationary Convolution

Weights

8 0 6

R

Inputs

W

Outputs

$$E = W \cdot \text{ceil}(R/2)^{\dagger}$$

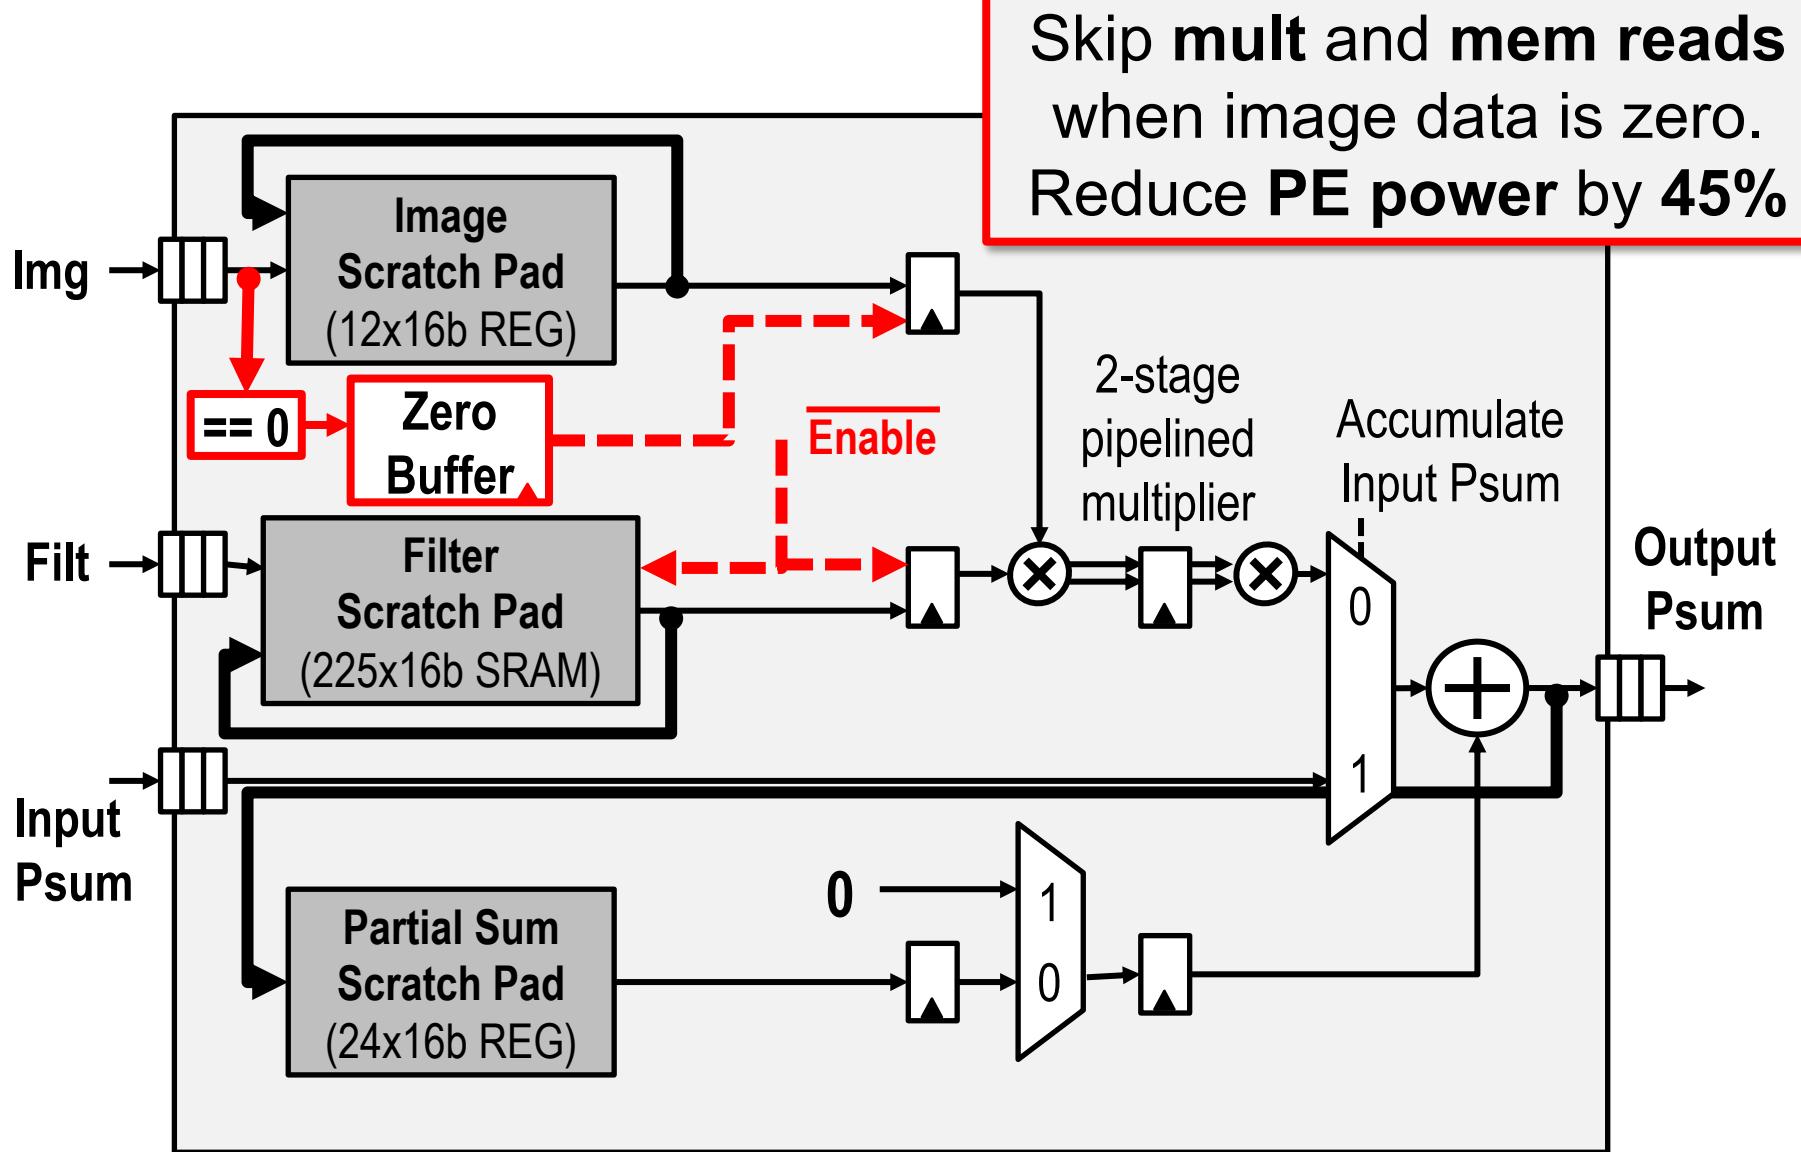
```
int i[W];      # Input activations
int w[R];      # Filter weights
int o[E];      # Output activations

for (e = 0; e < E; e++) {
    for (r = 0; r < R; r++) {
        if (!w[r]) o[e] += i[e+r]*w[r];
    }
}
```

Saved energy but not time

[†] Assuming: ‘valid’ style convolution

Eyeriss – Clock Gating



Compressed Weights

Compressed Storage of Weights

Uncompressed Weights

$$\begin{matrix} \text{Weights} \\ \boxed{8 \ 0 \ 6} \\ R \end{matrix} * \begin{matrix} \text{Inputs} \\ W \end{matrix} = \begin{matrix} \text{Outputs} \\ E = W \cdot \text{ceil}(R/2)^\dagger \end{matrix}$$

Compressed Weights

$$\text{Weights} \quad \text{Inputs} \quad \text{Outputs}$$
$$\begin{bmatrix} R_{11} & R_{12} \\ R_{21} & R_{22} \end{bmatrix} \times \begin{bmatrix} W_1 \\ W_2 \end{bmatrix} = \begin{bmatrix} E_1 \\ E_2 \end{bmatrix}$$
$$E = W \cdot \text{ceil}(R/2)^{\dagger}$$

[†] Assuming: ‘valid’ style convolution

Compressed Storage of Weights

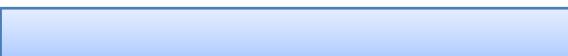
Uncompressed Weights

Weights

8	0	6
---	---	---

R

Inputs



=

Outputs

 $E = W \cdot \text{ceil}(R/2)^\dagger$

Compressed Weights

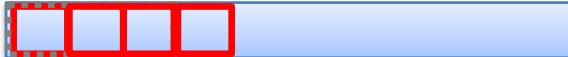
Weights

8	6
---	---

0	2
---	---

R

Inputs



=

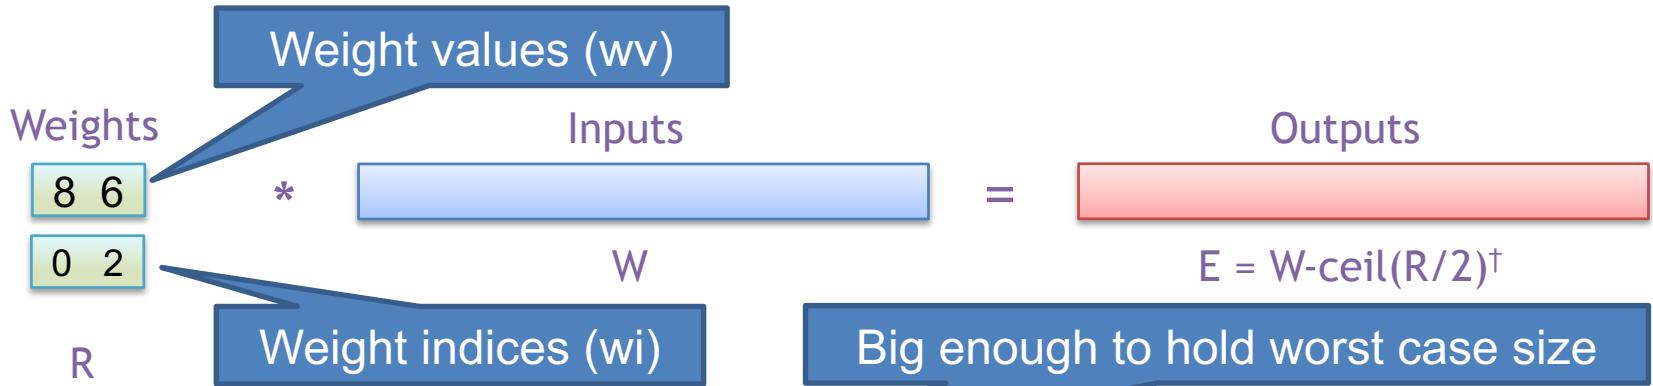
Outputs

 $E = W \cdot \text{ceil}(R/2)^\dagger$

Coordinate value in
uncompressed storage

[†] Assuming: ‘valid’ style convolution

Compressed Weights 1-D Convolution



```
int i[W];           # Input activations
int wv[R], wc[R]; # Compressed filter weights
int o[E];          # Output activations

for (e = 0; e < E; e++) {
    for (r = 0; r < R; r++) {
        break if !wi[r].valid;
        o[e] += i[e+wc[r]] * wv[r];
    }
}
```

Table lookup of weight index

No

Direct access into compressed buffer

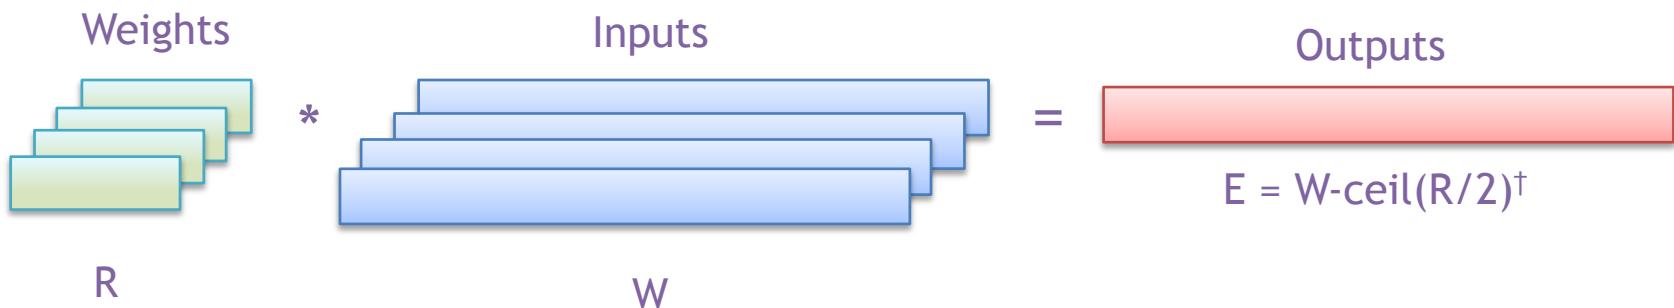
or pre-computed length

To Extend to Other Dimensions of DNN

- Need to add loop nests for:
 - 2-D input activations and filters
 - Multiple input channels
 - Multiple output channels
- Add parallelism...

Compressed Inputs

Multi-Input Channel 1-D Convolution



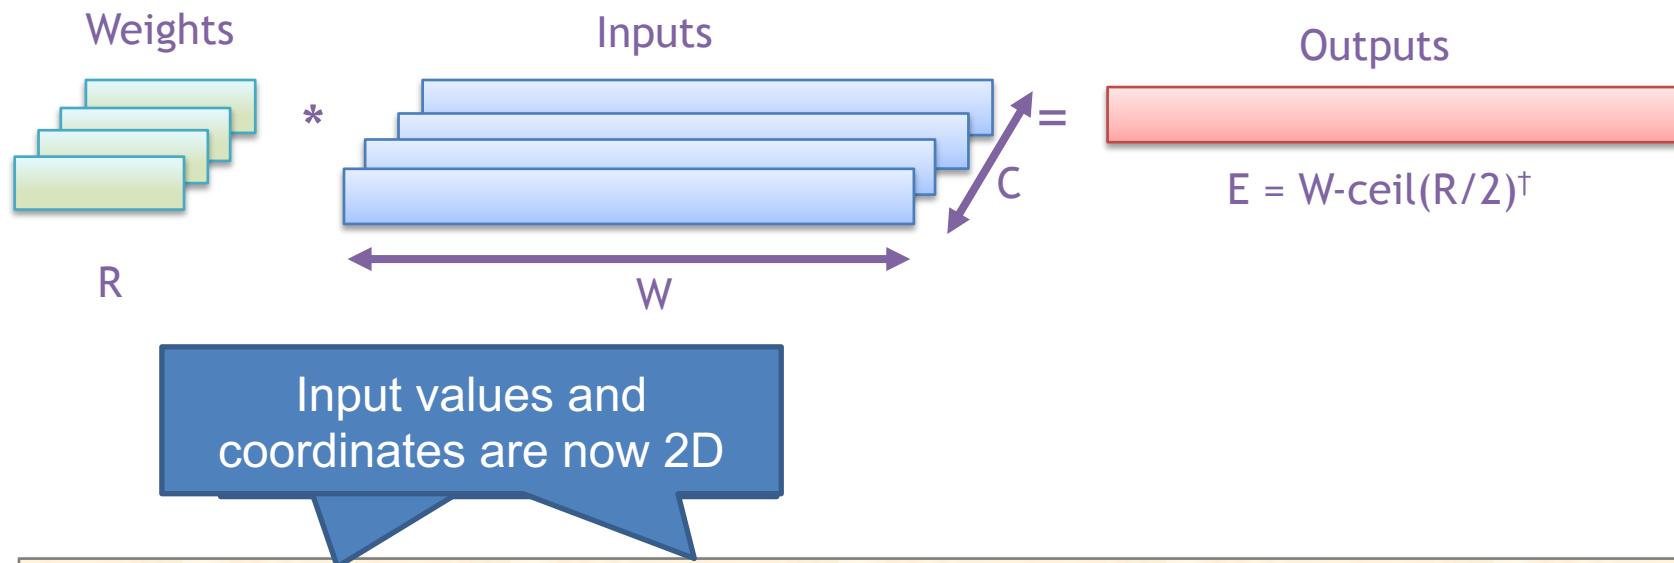
```
int i[C][W];           # Multi-input channel activations
int w[C][R];           # Filter weights
int o[E];              # Output activations

for (w = 0; w < W; w++) {
    for (r = 0; r < R; r++) {
        parallel-for (c = 0; c < C; c++) {
            o[w-r] += i[c][w]*w[c][r];
        }
    }
}
```

Note opportunity for
spatial sum

[†] Assuming: ‘valid’ style convolution

Multi-Input Channel 1-D Convolution

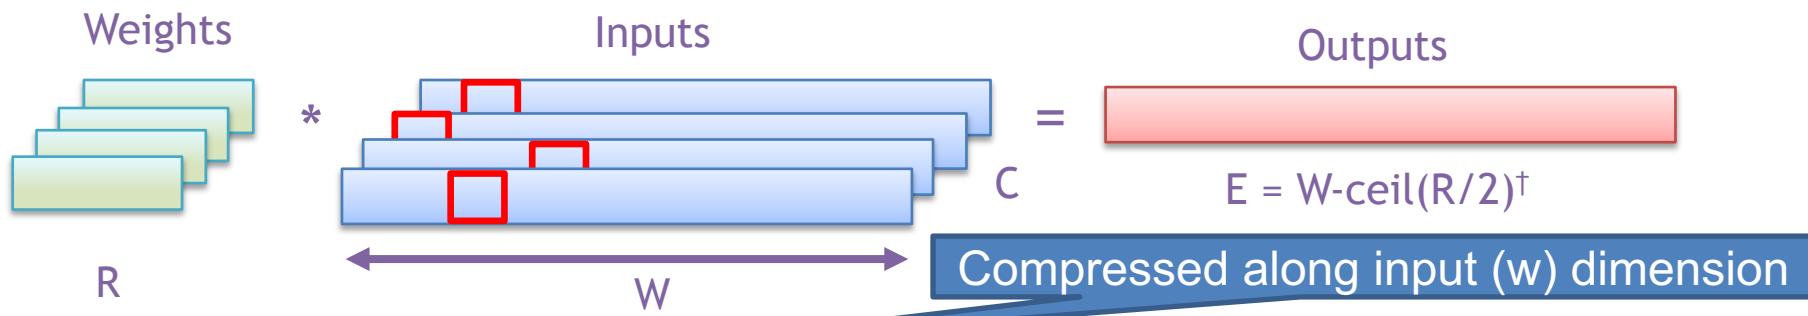


```
int iv[C][W], ic[C][W];      # Compressed input activations  
int w[C][R];                 # Filter weights  
int o[E];                     # Output activations
```

Should we compress along C or W dimension? **Let's see**

[†] Assuming: ‘valid’ style convolution

Compressed Sparse W-dimension



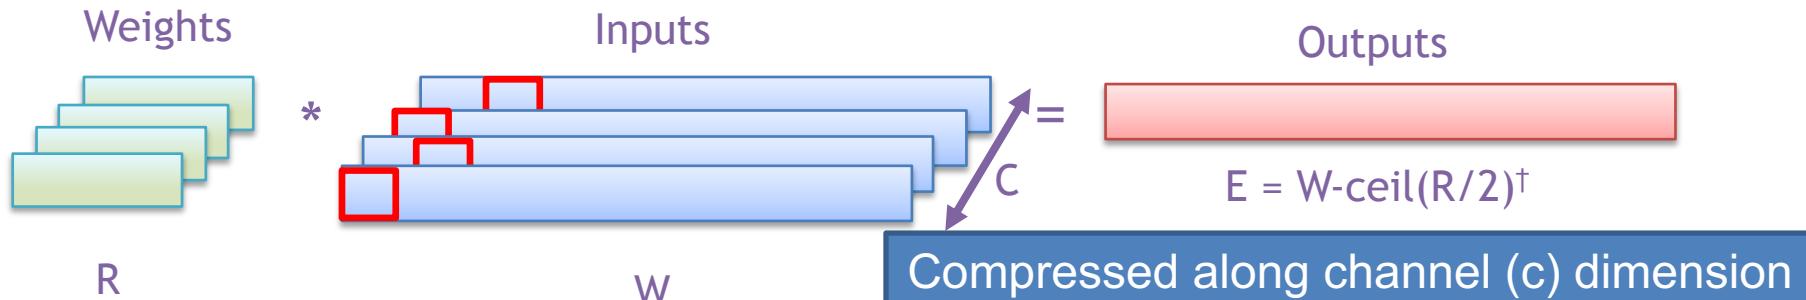
```
int iv[C][W], icw[C][W]; # Compressed input activations  
int w[C][R]; # Filter weights  
int o[E]; # Output activations
```

```
for (w = 0; w < W; w++) {  
    for (r = 0; r < R; r++) {  
        parallel-for (c = 0; c < C; c++) {  
            break if !icw[c][w].valid;  
            o[icw[c][w]-r] += iv[c][w]*w[c][r];  
    }}}
```

Running parallel COMPRESSed w's

The variation of these index values with different c's will prevent synchronized spatial sum

Compressed Sparse C-dimension



```
int iv[C][W], icc[C][W]; # Compressed input activations  
int w[C][R]; # Filter weights  
int o[E]; # Output activations
```

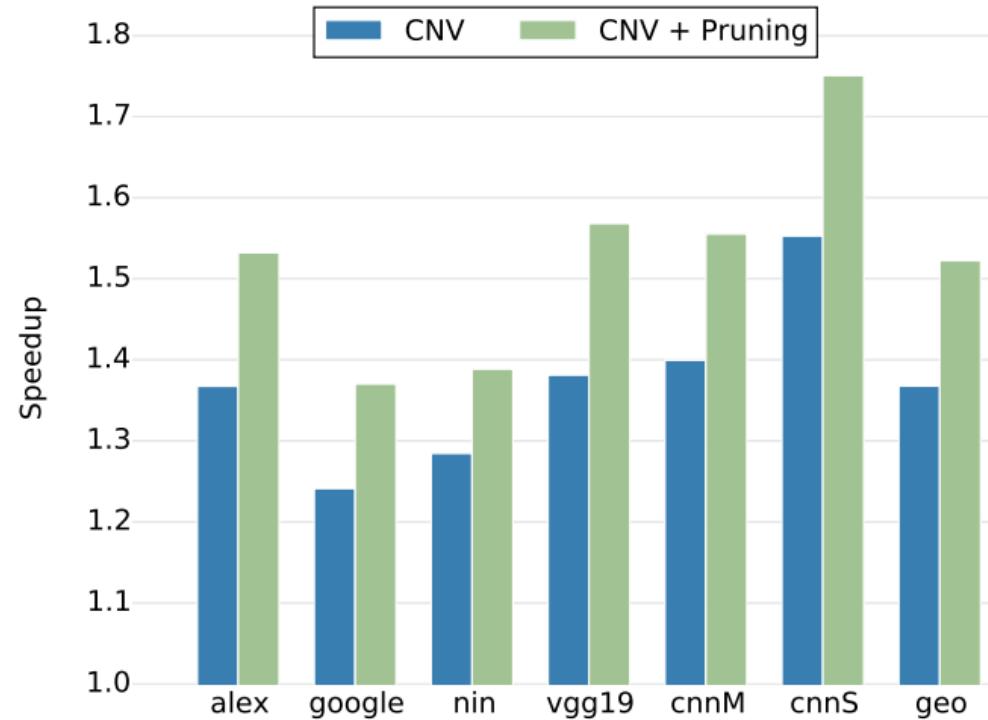
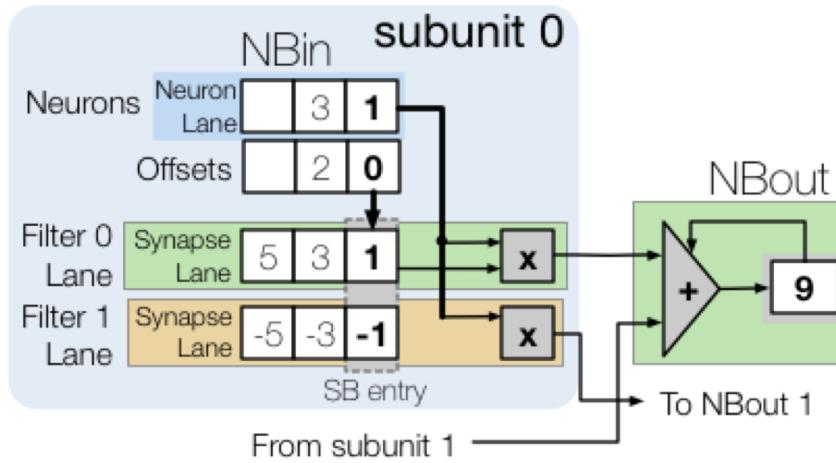
```
for (r = 0; r < R; r++) {  
    for (w = 0; w < W; w++) {  
        parallel-for (c = 0; c < C; c++) {  
            break if !icc[c][w].valid;  
            o[w-r] += iv[c][w]*w[icc[c][w]][r];  
    }}}
```

Running parallel COMPRESSed c's

Note we now have a synchronized spatial sum

Cnvlutin

- Process Convolution Layers
- Built on top of DaDianNao (4.49% area overhead)
- Speed up of 1.37x (1.52x with activation pruning)



Compressing Inputs + Weights

Output Stationary – Sparse W&I

Weights	Inputs	Outputs
8 0 6	4 0 0 0 3 0 0 8 0 2	1/3 0/3 1/3 0/3
R	W	E = W-ceil(R/2) [†]

```
int i[W];                      # Input activations
int w[R];                      # Filter weights
int o[E];                      # Output activations

for (e = 0; e < E; e++) {
    parallel-for (r = 0; r < R; r++) {
        next if w[r] == 0;
        next if i[e+r] == 0;
        o[e] += i[e+r] * w[r];
    }
}
```

How often is work done in inner loop?

Not very much!

Flattened Inputs & Weights

```
int i[C][W*H];           # Flattened input activations
int w[C][M*R*S];         # Flattened filter weights
int o[M][E][F];          # Output activations

for mrs2 = [0..MRS2) {
    for c2 = [0..C) {
        for wh1 = [0..WH1) {
            for mrs1 = [0..MRS1) {
                parallel-for wh0 = [0..WH0) x
                    mrs0 = [0..MRS0) {
                        m = Mcoord(mrs2, mrs1, mrs0);
                        e = Wcoord(wh1, wh0) - Rcoord(mrs2, mrs1, mrs0);
                        f = Hcoord(wh1, wh0) - Scoord(mrs2, mrs1, mrs0);
                        o[m][e][f] += i[c2][wh1*WH0+wh0]
                                      * w[c2][mrs2*MRS1*MRS0+mrs1*MRS0+mrs0];
                    }
            }
        }
    }
}
```

At inner loop inputs are stationary across steps mrs1

Any opportunity for spatial sum? No

Sparse CNN (SCNN)

- Architecture to exploit sparsity

Intuition behind SCNN

Forget the sliding windows based convolution

Observation

Each non-zero activation must be multiplied by each non-zero weight

Intuition behind SCNN

Forget the sliding windows based convolution

Observation

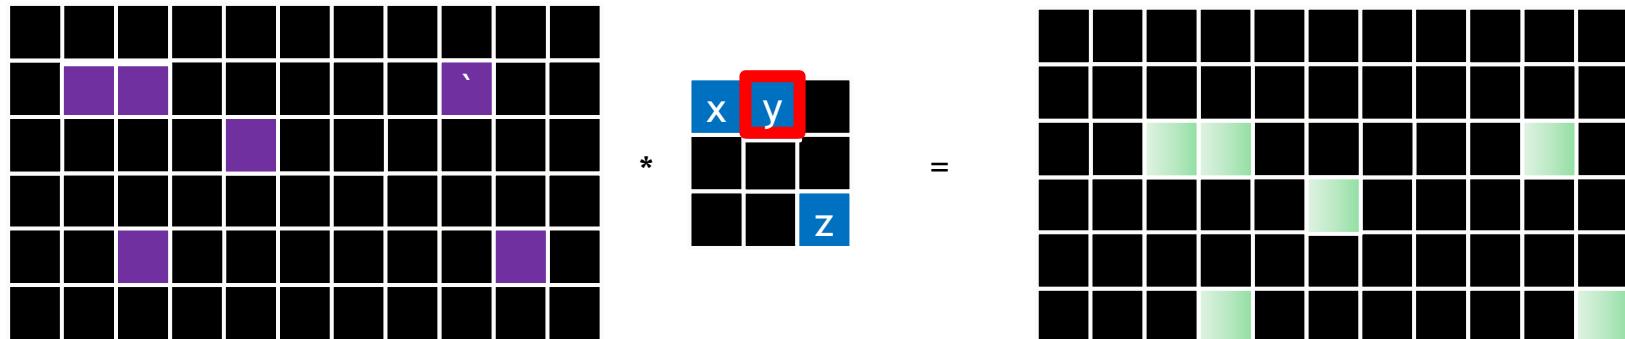
Each non-zero activation must be multiplied by each non-zero weight

Intuition behind SCNN

Forget the sliding windows based convolution

Observation

Each non-zero activation must be multiplied by each non-zero weight

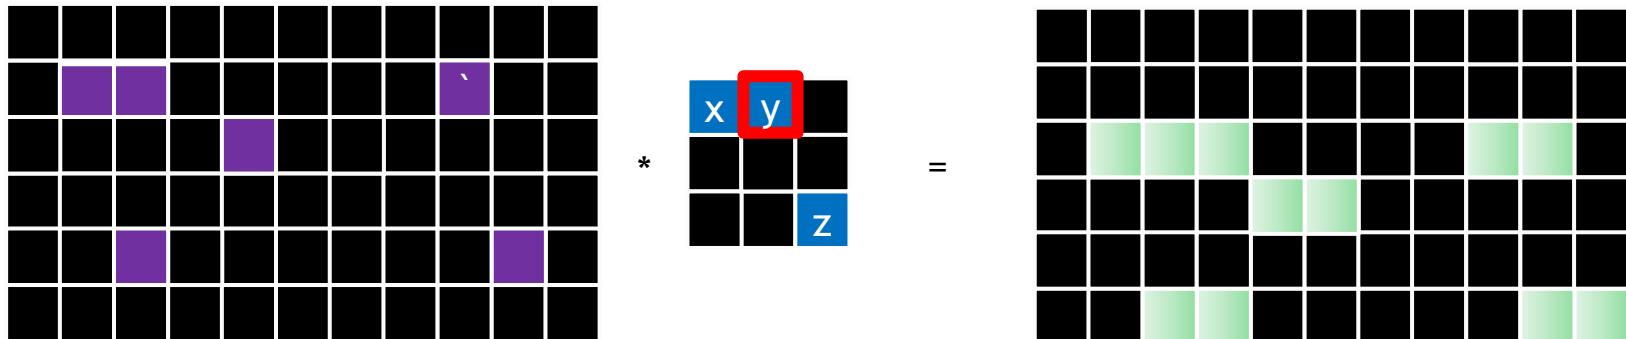


Intuition behind SCNN

Forget the sliding windows based convolution

Observation

Each non-zero activation must be multiplied by each non-zero weight



Intuition behind SCNN

Forget the sliding windows based convolution

Observation

Each non-zero activation must be multiplied by each non-zero weight

$$\begin{matrix} a & b \\ & d \\ e & \end{matrix} * \begin{matrix} x & y \\ & z \end{matrix} = ?$$

A diagram illustrating the computation of a sum of products (SCNN) from two input matrices. The first matrix has non-zero elements labeled 'a', 'b', 'd', 'e', and 'f' in red boxes. The second matrix has non-zero elements labeled 'x', 'y', and 'z' in blue boxes. The result is a matrix where the value at position (3,3) is a question mark, indicating that the product of 'e' and 'z' is being calculated.

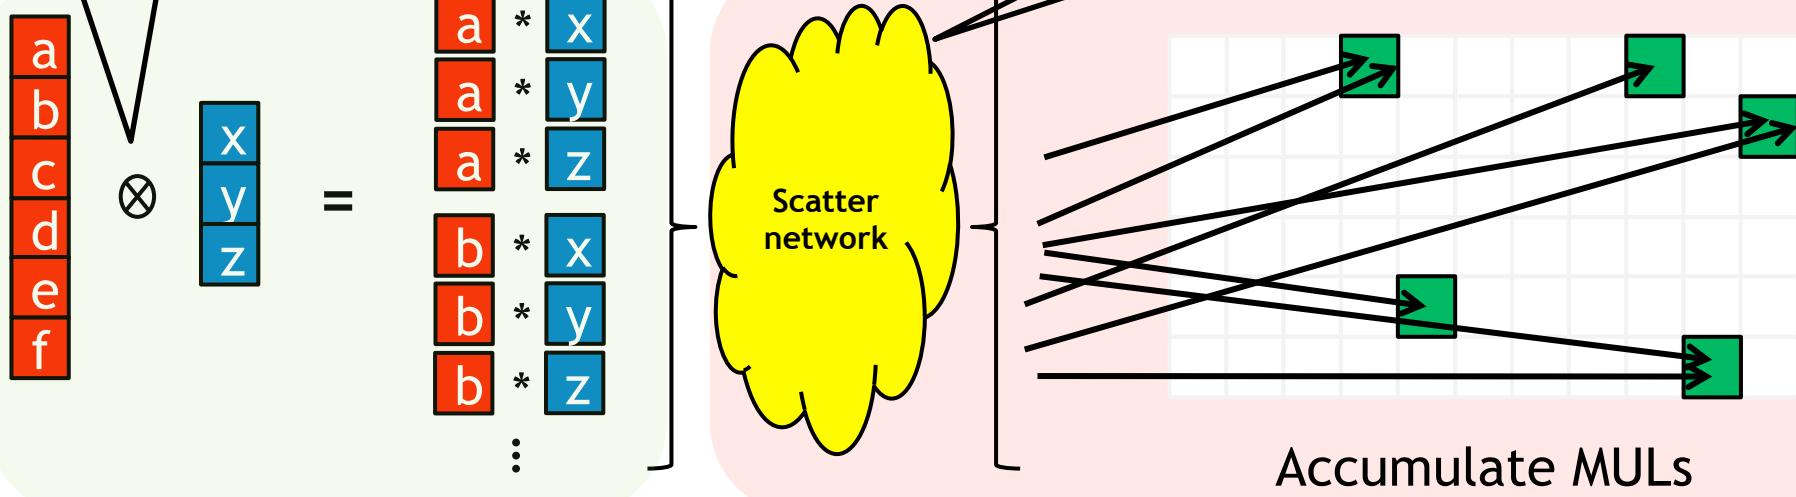
Sparse CNN (SCNN)

Supports Convolutional Layers

Densely Packed Storage of Weights and Activations

All-to all Multiplication of Weights and Activations

Mechanism to Add to Scattered Partial Sums



PE frontend

PE backend

Input Stationary Dataflow

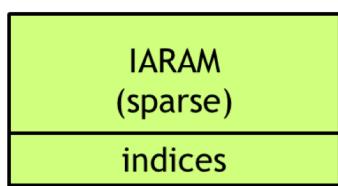
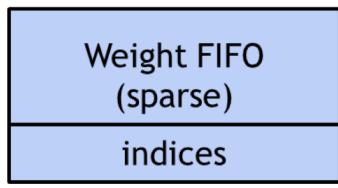
[Parashar et al., ISCA 2017]

SCNN PE microarchitecture

Sparse-compressed frontend

Flattened Weights

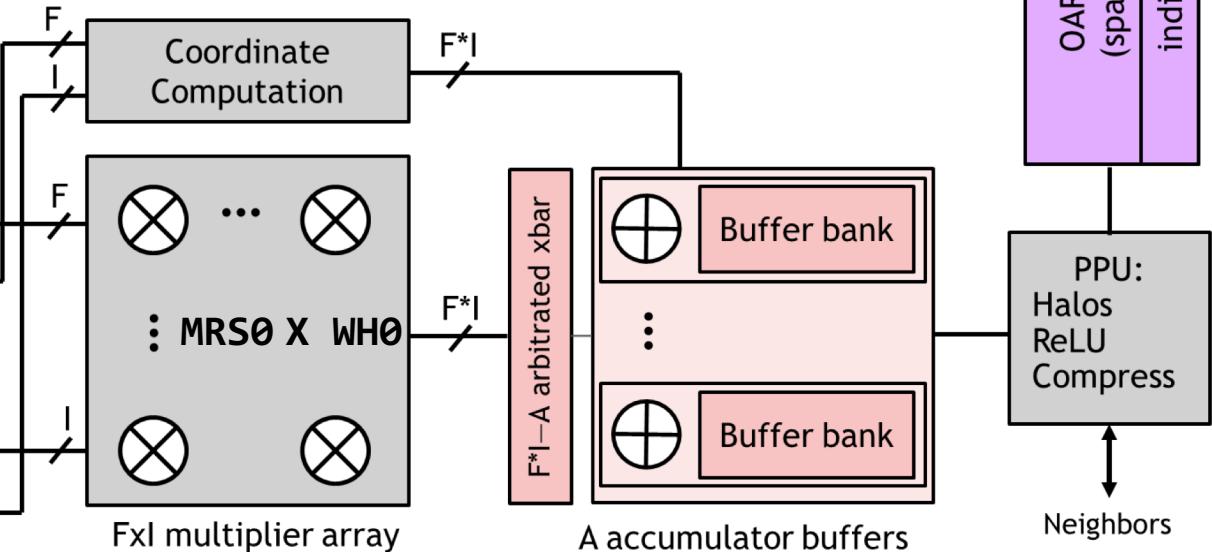
- wv[C][M*R*S],
- wim[C][M*R*S],
- wir[C][M*R*S],
- wis[C][M*R*S];



Flattened Input Activations

- iv[C][W*H]
- iiw[C][W*H],
- iuh[C][W*H];

$m = \text{Mcoord}(\text{mrs}_2, \text{mrs}_1, \text{mrs}_0);$
 $e = \text{Wcoord}(\text{wh}_1, \text{wh}_0) - \text{Rcoord}(\text{mrs}_2, \text{mrs}_1, \text{mrs}_0);$
 $f = \text{Hcoord}(\text{wh}_1, \text{wh}_0) - \text{Scoord}(\text{mrs}_2, \text{mrs}_1, \text{mrs}_0);$



Dense backend

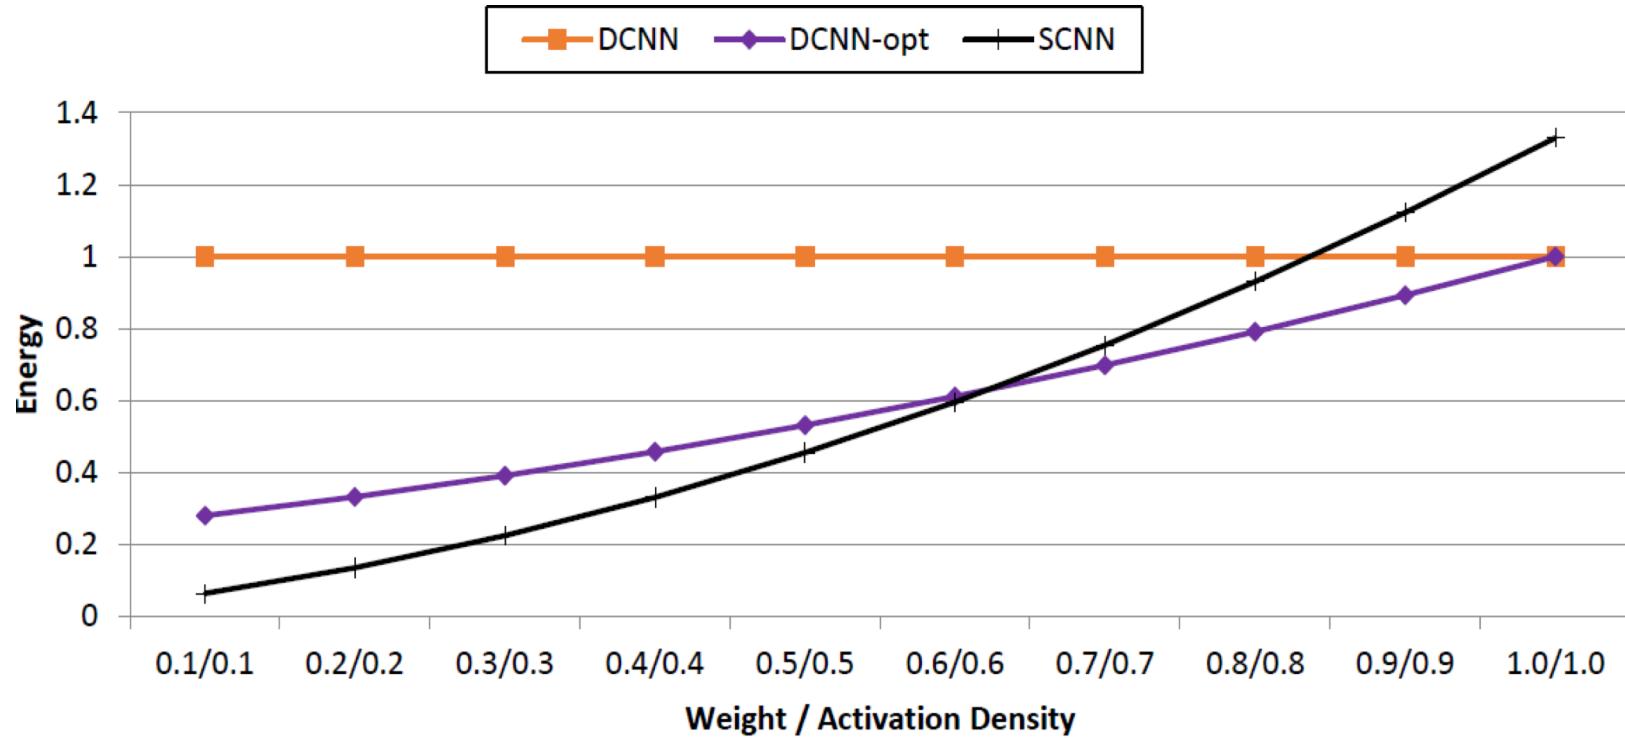
[Parashar et al., ISCA 2017]

Flattened Inputs & Weights

```
int iv[C][W*H] iiw[C][W*H], iih[C][W*H];
int wv[C][M*R*S], wim[C][M*R*S], wir[C][M*R*S], wis[C][M*R*S];
int o[M][E][F];

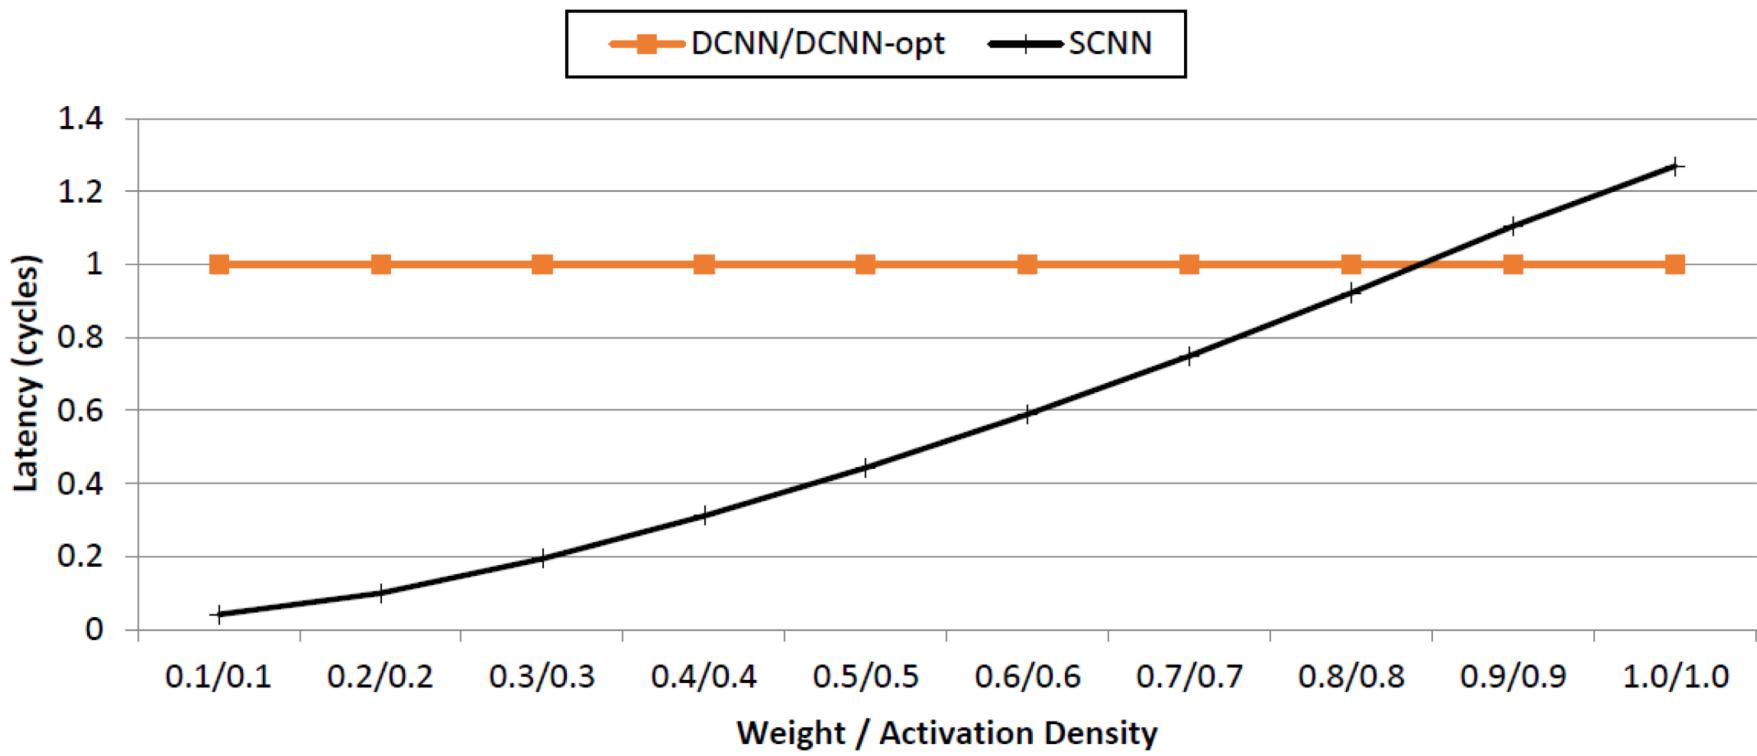
for mrs2 = [0..MRS2) {
    for c2 = [0..C) {
        for wh1 = [0..WH1) {
            for mrs1 = [0..MRS1) {
                parallel-for wh0 = [0..WH0) x
                    mrs0 = [0..MRS0) {
                        break if !ii[c2][mrs2*MRS1*MRS0+mrs1*MRS0+mrs0].v;
                        break if !wi[c2][wh1*WH0+wh0].v;
                        m = Mcoord(mrs2, mrs1, mrs0);
                        e = Wcoord(wh1, wh0)-Rcoord(mrs2, mrs1, mrs0);
                        f = Hcoord(wh1, wh0)-Scoord(mrs2, mrs1, mrs0);
                        o[m][e][f] += i[c2][wh1*WH0+wh0]
                                    * w[c2][mrs2*MRS1*MRS0+mrs1*MRS0+mrs0];
                }
            }
        }
    }
}
```

SCNN Energy Versus Density



[Parashar et al., ISCA 2017]

SCNN Latency Versus Density



[Parashar et al., ISCA 2017]

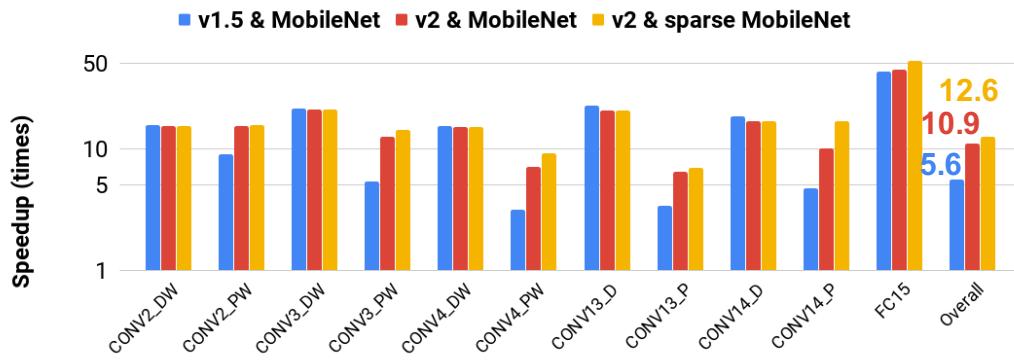
Eyeriss – V2

- Architecture to accommodate variety sparsity

Eyeriss v2: Balancing Flexibility and Efficiency

Efficiently supports

- Wide range of filter shapes
 - Large and Compact
- Different Layers
 - CONV, FC, depth wise, etc.
- **Wide range of sparsity**
 - Dense and Sparse
- Scalable architecture



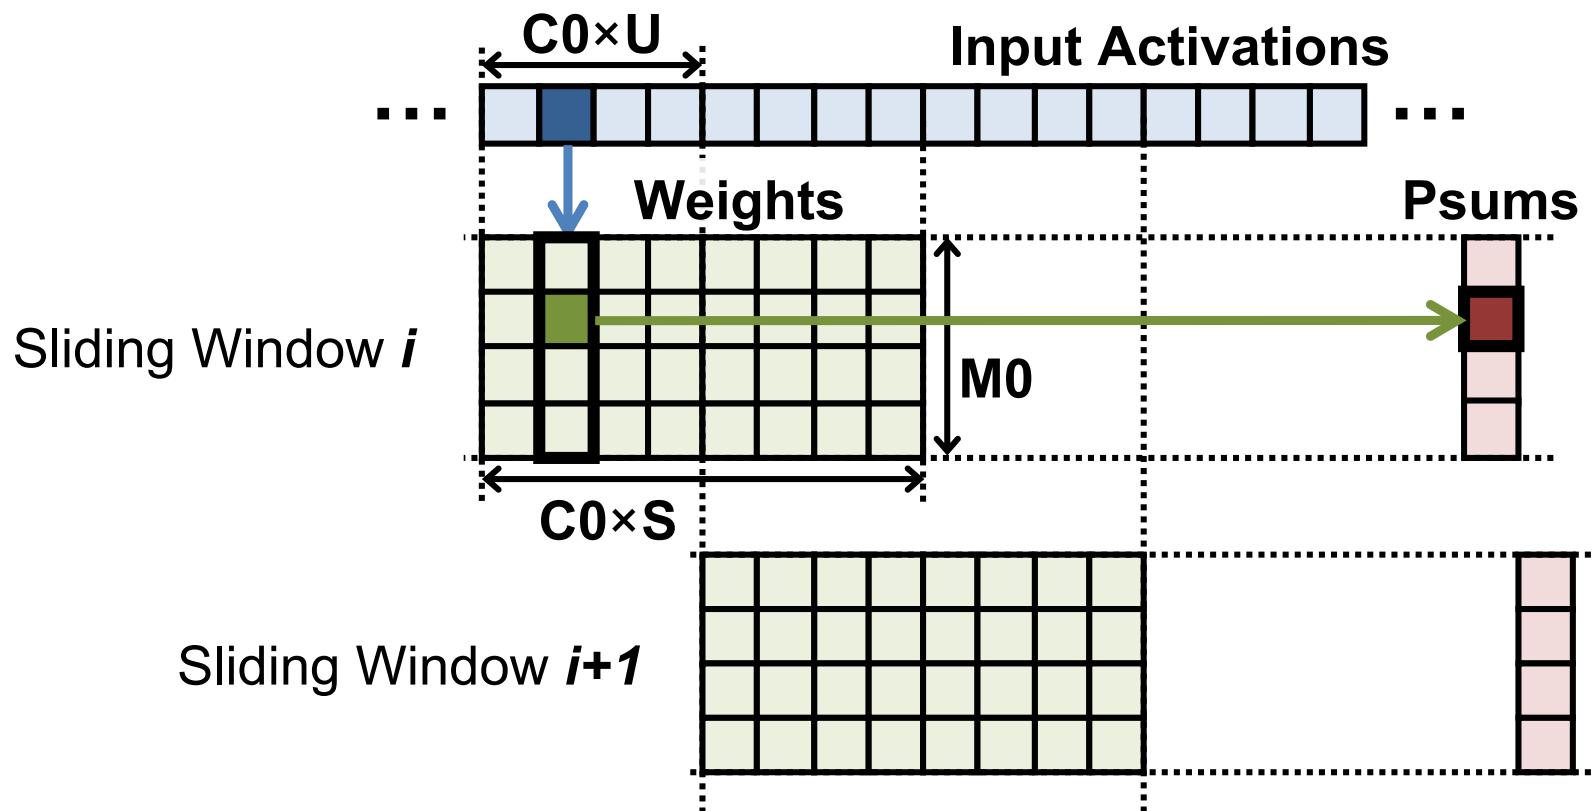
Speed up over Eyeriss v1 scales with number of PEs

# of PEs	256	1024	16384
AlexNet	17.9x	71.5x	1086.7x
GoogLeNet	10.4x	37.8x	448.8x
MobileNet	15.7x	57.9x	873.0x

[Chen et al., JETCAS 2019]

Over an order of magnitude faster and more energy efficient than Eyeriss v1

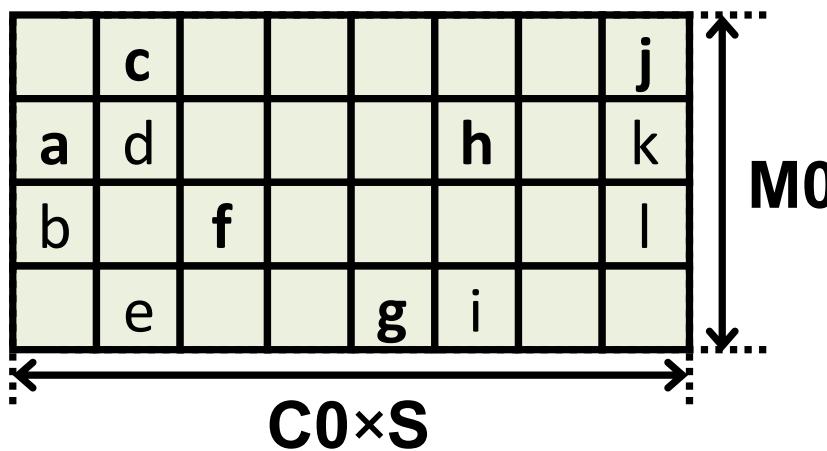
Eyeriss v2: Processing In PE



- **M0** : # of output channels processed in a PE
- **C0** : # of input channels processed in a PE
- **S** : filter width
- **U** : stride

Eyeriss v2: Compressed Data Format

Weight Matrix



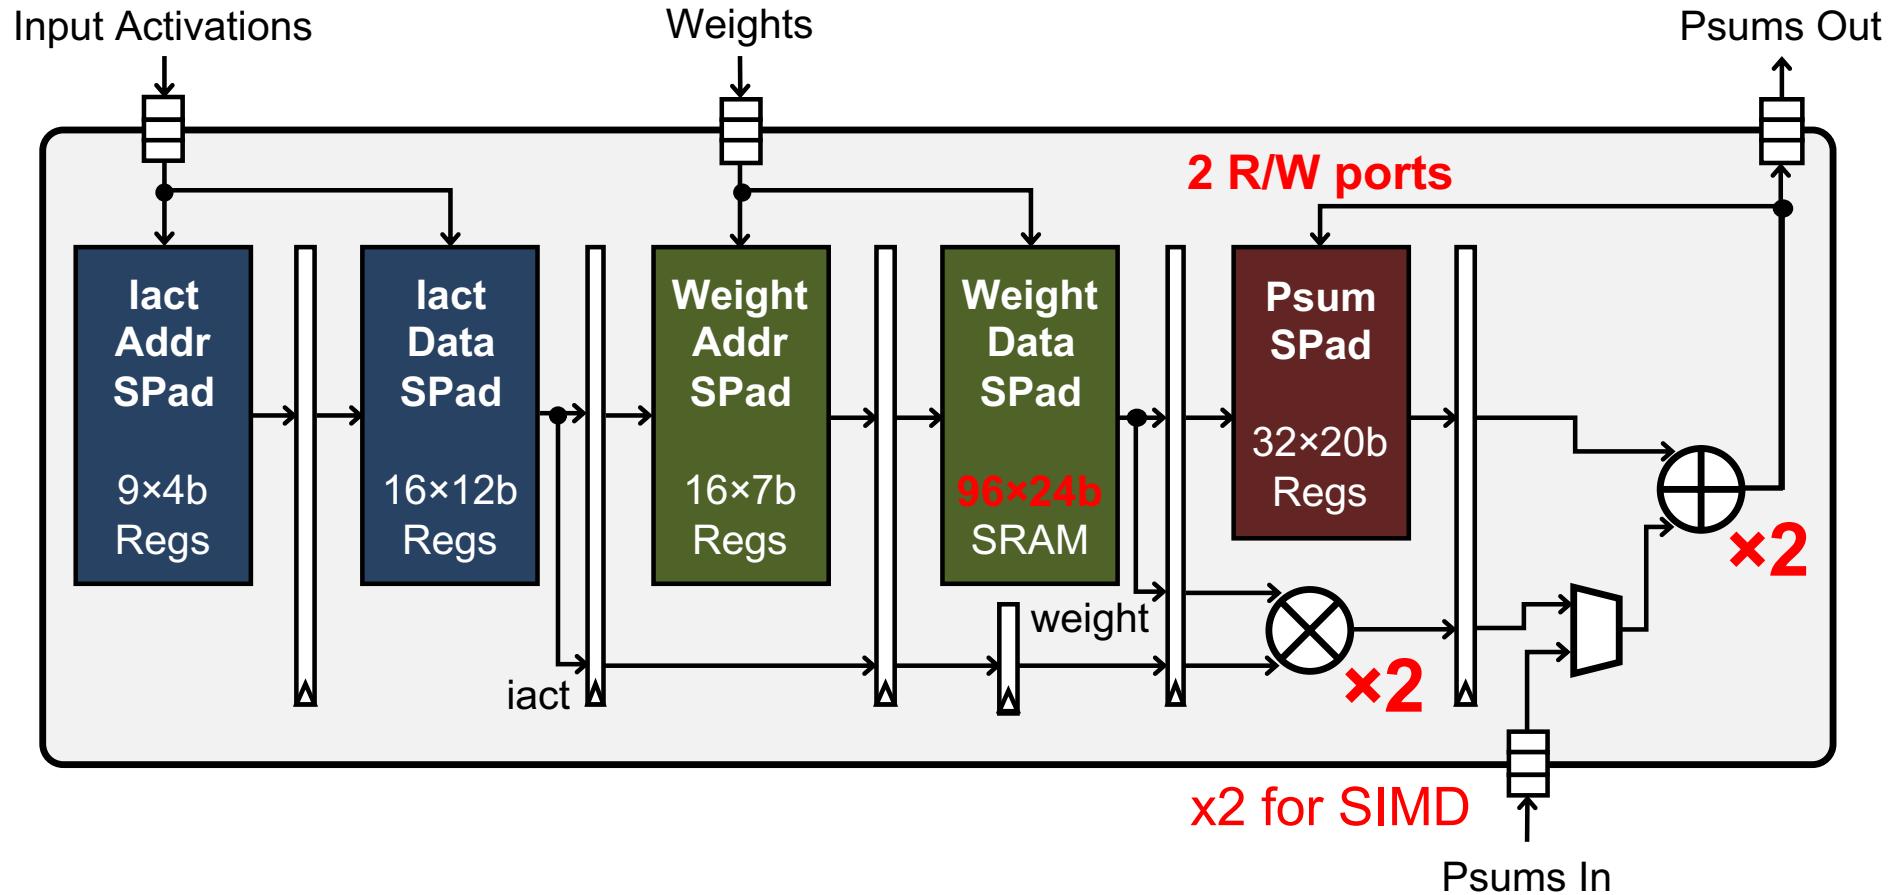
CSC Compressed Data:

data vector: {**a, b, c, d, e, f, g, h, i, j, k, l**}

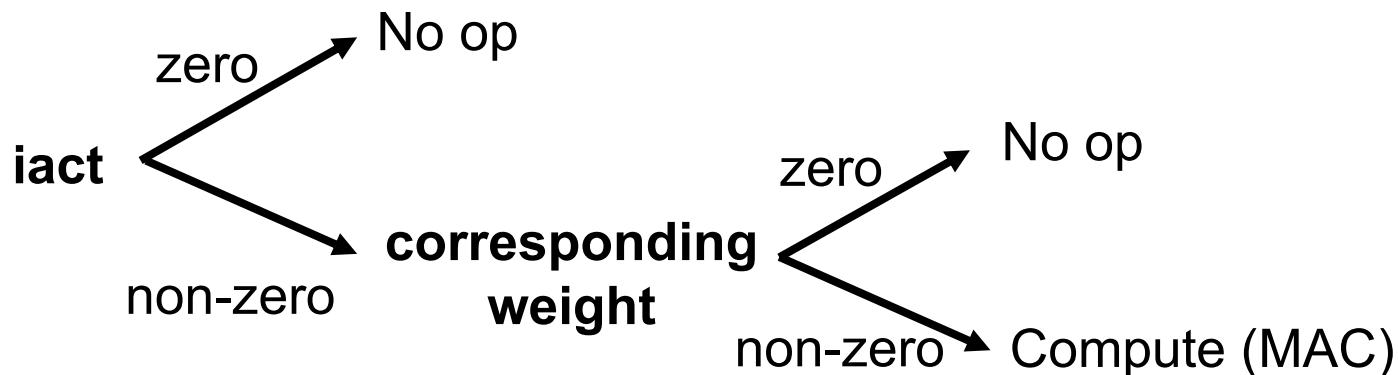
count vector: {1, 0, 0, 0, 1, 2, 3, 1, 1, 0, 0, 0}

address vector: {0, 2, 5, 6, 6, 7, 9, 9, 12}

Eyeriss v2: PE Architecture



Decision Tree in Eyeriss v2 PE



- **If the iact is zero**, the CSC format will ensure that it is not read from the spad and therefore no cycles are wasted.
- **If the iact is not zero**, its value will be fetched from the iact data SPad and passed to the next pipeline stage.
 - **If there are non-zero weights corresponding to the non-zero iacts**, they will be passed down the pipeline for computation. The zero weights will be skipped since the weights are also encoded with the CSC format.
 - **If there are no non-zero weights corresponding to the non-zero iacts**, the non-zero iacts will not be further passed down in the pipeline.

Summary

- Processing Irregular (Gather-Scatter)
 - If weights and inputs compressed to dense (gather); output scatter
 - If weights and inputs uncompressed sparse (scatter); output gather
- Overhead (must not exceed benefits of sparsity)
 - Storage of location information for compressed data
 - Logic for checking if either inputs are zero
- Underutilization
 - Number of parallel cores (tiling) → maximize parallelism, but minimize underutilization
 - Flatten to 1-D avoid fragmentation from limits of each dimension
- Workload Imbalance
- **Lots of challenges in sparse deep neural network acceleration!**